

Yuntian DENG

yuntiandeng.com • yuntian@uwaterloo.ca • Google Scholar

Employment

University of Waterloo	Waterloo, Canada
Assistant Professor, Cheriton School of Computer Science	Aug 2024–Present
Vector Institute	Toronto, Canada
Faculty Affiliate	Aug 2024–Present
Harvard University	Remote
Associate in Computer Science	Jun 2024–Present
NVIDIA	Toronto, Canada
Visiting Professor (manager: Yejin Choi)	Nov 2024–Nov 2025
Allen Institute for Artificial Intelligence	Seattle, WA
Postdoc (advised by Yejin Choi)	Jul 2023–Jul 2024

Research Interests

ProgramAsWeights (PAW) and neural programs; neural world models and interactive systems; implicit chain-of-thought reasoning and internalized computation; real-world LLM use and open research infrastructure.

Education

Harvard University	Cambridge, MA
Ph.D. in Computer Science	May 2023
Advisors: Alexander Rush and Stuart Shieber	
Carnegie Mellon University	Pittsburgh, PA
Master of Language Technologies	Aug 2016
Advisor: Eric Xing	
Tsinghua University	Beijing, China
Bachelor of Engineering, Department of Automation	Jul 2014

Awards

Laude Institute Slingshots (NeuralOS)	Jun 2026
Google Research Award (co-PI)	Apr 2026
Argonne National Lab Impact Award	Jan 2023
University of Chicago Rising Stars in Data Science	Nov 2022
ACM Gordon Bell Special Prize for COVID-19 Research	Nov 2022
NVIDIA Graduate Fellowship	Dec 2021
Microsoft Turing Academic Program (implicit chain-of-thought research)	Dec 2021
Harvard Certificates of Distinction in Teaching	Spring 2019, Fall 2020, Fall 2021
Twitch Fellowship Finalist	Jan 2021
DAC 2020 Best Paper Award	Jul 2020
Baidu Fellowship	2019
French American Doctoral Exchange Program (French Embassy, 10 US laureates)	Jun 2018
ACL 2017 Best Demonstration Paper Runner-Up	Aug 2017

Research Contributions

ProgramAsWeights (PAW). Compiles English function specifications into compact neural programs that run locally; its public hub contains 3,800+ community-built programs.

Representative work: Program-as-Weights: A Programming Paradigm for Fuzzy Functions (*arXiv 2026*)

Project: programasweights.com

NeuralOS. Uses a neural generative model to simulate an interactive operating system end to end; its public deployment has 50,000+ visits.

Representative work: NeuralOS: Towards Simulating Operating Systems via Neural Generative Models (*ICLR 2026*)

Project: neural-os.com

Implicit Chain of Thought. Introduced implicit chain-of-thought reasoning as a research direction and helped catalyze its growth, including by co-organizing the Latent & Implicit Thinking Workshop at ICLR 2026, which accepted 78 papers.

Representative works: Implicit Chain of Thought Reasoning via Knowledge Distillation (*arXiv 2023*); From Explicit CoT to Implicit CoT (*arXiv 2024*)

Demo: A 12-layer GPT-2 achieves 99.5% accuracy when multiplying two 20-digit numbers without emitting intermediate reasoning

WildChat. Provides 4.8 million real user–ChatGPT conversations and an open visualizer; it is among the AI Observatory’s largest and most detailed datasets and is used by OpenAI and Anthropic for safety evaluation.

Representative work: WildChat: 1M ChatGPT Interaction Logs in the Wild (*ICLR 2024 Spotlight*)

Dataset: WildChat-4.8M-Full on Hugging Face; *Visualizer:* wildvisualizer.com

Selected Media Coverage

Interviews and Direct Features

The Canadian Press: Researchers develop tool to help courts spot AI-generated evidence	Dec 5, 2025
NZZ: What people really ask ChatGPT (English translation)	Nov 30, 2024
TechCrunch: Why is ChatGPT so bad at math?	Oct 2, 2024
The Washington Post: What do people really ask chatbots?	Aug 4, 2024

Coverage of Coauthored Research

The Washington Post homepage (The 7): New data shows people are getting personal with AI	Aug 18, 2026
MIT Technology Review: We still don’t know how people are really using AI	Aug 18, 2026

Publications and Preprints

* Equal contribution; † Equal advising.

- 50 **Yuntian Deng**, Pengyu Nie, Stuart Shieber. Compile by Training: Turning Natural-Language Specifications into Local Neural Functions. *EMNLP 2026 Demo*
- 49 Wentao Zhang, Shariyar Murtaza, Junaid Bhatti, Utkarsh Soni, Yifan Nie, Eugene Wen, **Yuntian Deng**. Control-Data Flow Separation: Stable Prompt Optimization in Multi-Agent LLMs. *EMNLP 2026 Findings*
- 48 Wentao Zhang*, Xuanhe Pan*, Han Zhou*, Yang Lu†, **Yuntian Deng**†. Interactive Training 2: Auditable Control Plane for Live Model Training. *arXiv 2026*
- 47 Wentao Zhang*, Liliana Hotsko*, Woojeong Kim*, Pengyu Nie, Stuart Shieber, **Yuntian Deng**. Program-as-Weights: A Programming Paradigm for Fuzzy Functions. *arXiv 2026*
- 46 Kelly McConvey, Jalehsadat Mahdavi-moghaddam, Nima Jamali, Maksym Taranukhin, Sajad Ebrahimi, Wentao Zhang, Karen Eltis, **Yuntian Deng**, Maura R. Grossman, Vered Shwartz, Ebrahim Bagheri. The CIFAR Synthetic Evidence Corpus for Detecting AI-Generated Evidence. *arXiv 2026*

- 45 Liliana Hotsko, Yinxi Li, **Yuntian Deng**, Pengyu Nie. Code2LoRA: Hypernetwork-Generated Adapters for Code Language Models under Software Evolution. *arXiv 2026*
- 44 Sajad Ebrahimi, Nima Jamali, Bardia Shirsalimian, Kelly McConvey, Wentao Zhang, Jalehsadat Mahdavi Moghaddam, Maksym Taranukhin, Maura R. Grossman, Vered Shwartz, **Yuntian Deng**, Ebrahim Bagheri. DetectZoo: A Unified Toolkit for AI-Generated Content Detection Across Text, Audio, and Image Modalities. *arXiv 2026*
- 43 Shayne Longpre, Anka Reuel, Dayeon Ki, Zhiping Zhang, Cathy Mengying Fang, Chuanyang Jin, Jennifer Mickel, Cedric Deslandes Whitney, Niloofar Mireshghallah, Victor Ojewale, Megan Richards, Ariel N. Lee, Alex Pentland, Tianshi Li, **Yuntian Deng**, Sara Hooker, Mykel Kochenderfer, Sanmi Koyejo. The AI Observatory: A Public Measure of Real-World AI Use. *Preprint*
- 42 David Acuna, Chao-Han Huck Yang, **Yuntian Deng**, Jaehun Jung, Ximing Lu, Prithviraj Ammanabrolu, Hyunwoo Kim, Yuan-Hong Liao, Yejin Choi. Long Grounded Thoughts: Synthesizing Visual Problems and Reasoning Chains at Scale. *ICML 2026*
- 41 Yinxi Li, **Yuntian Deng**, Pengyu Nie. TokDrift: When LLM Speaks in Subwords but Code Speaks in Grammar. *ACL 2026*
- 40 Luke Rivard, Sun Sun, Hongyu Guo, Wenhui Chen, **Yuntian Deng**. NeuralOS: Towards Simulating Operating Systems via Neural Generative Models. *ICLR 2026*
- 39 Wei Liu, Ruochen Zhou, Yiyun Deng, Yuzhen Huang, Junteng Liu, **Yuntian Deng**, Yizhe Zhang, Junxian He. Learn to Reason Efficiently with Adaptive Length-Based Reward Shaping. *ICLR 2026*
- 38 Shivalika Singh, Yiyang Nan, Alex Wang, Daniel D’Souza, Sayash Kapoor, Ahmet Üstün, Sanmi Koyejo, **Yuntian Deng**, Shayne Longpre, Noah A. Smith, Beyza Ermis, Marzieh Fadaee, Sara Hooker. The Leaderboard Illusion. *NeurIPS 2025 Datasets & Benchmarks Track*
- 37 Wentao Zhang, Yang Young Lu, **Yuntian Deng**. Interactive Training: Feedback-Driven Neural Network Optimization. *EMNLP 2025 Demo*
- 36 Wentao Zhang, Woojeong Kim, **Yuntian Deng**. From Chat Logs to Collective Insights: Aggregative Question Answering. *EMNLP 2025 Oral*
- 35 Xiaoyan Bai, Itamar Pres, **Yuntian Deng**, Chenhao Tan, Stuart Shieber, Fernanda Viégas, Martin Wattenberg, Andrew Lee. Why Can’t Transformers Learn Multiplication? Reverse-Engineering Reveals Long-Range Dependency Pitfalls. *arXiv 2025*
- 34 Yifan Zong, **Yuntian Deng**, Pengyu Nie. Mix-of-Language-Experts Architecture for Multilingual Programming. *LLM4Code 2025*
- 33 Bill Yuchen Lin, **Yuntian Deng**, Khyathi Chandu, Faeze Brahman, Abhilasha Ravichander, Valentina Pyatkin, Nouha Dziri, Ronan Le Bras, Yejin Choi. WildBench: Benchmarking LLMs with Challenging Tasks from Real Users in the Wild. *ICLR 2025 Spotlight*
- 32 Zhangchen Xu, Fengqing Jiang, Luyao Niu, **Yuntian Deng**, Radha Poovendran, Yejin Choi, Bill Yuchen Lin. Magpie: Alignment Data Synthesis from Scratch by Prompting Aligned LLMs with Nothing. *ICLR 2025*
- 31 Jinjie Ni, Fuzhao Xue, Xiang Yue, **Yuntian Deng**, Mahir Shah, Kabir Jain, Graham Neubig, Yang You. MixEval: Deriving Wisdom of the Crowd from LLM Benchmark Mixtures. *NeurIPS 2024*
- 30 **Yuntian Deng**, Wenting Zhao, Jack Hessel, Xiang Ren, Claire Cardie, Yejin Choi. WildVis: Open Source Visualizer for Million-Scale Chat Logs in the Wild. *EMNLP 2024 Demo*
- 29 Wenting Zhao, Tanya Goyal, Yu Ying Chiu, Liwei Jiang, Benjamin Newman, Abhilasha Ravichander, Khyathi Chandu, Ronan Le Bras, Claire Cardie, **Yuntian Deng**, Yejin Choi. WildHallucinations: Evaluating Long-form Factuality in LLMs with Real-World Entity Queries. *arXiv 2024*
- 28 **Yuntian Deng**, Yejin Choi, Stuart Shieber. From Explicit CoT to Implicit CoT: Learning to Internalize CoT Step by Step. *arXiv 2024*

- 27 Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, **Yuntian Deng**. WildChat: 1M ChatGPT Interaction Logs in the Wild. *ICLR 2024 Spotlight*
- 26 John X. Morris, Chandan Singh, Alexander Rush, Jianfeng Gao, **Yuntian Deng**. Tree Prompting: Efficient Task Adaptation without Fine-Tuning. *EMNLP 2023*
- 25 **Yuntian Deng**, Kiran Prasad, Roland Fernandez, Paul Smolensky, Vishrav Chaudhary, Stuart Shieber. Implicit Chain of Thought Reasoning via Knowledge Distillation. *arXiv 2023*
- 24 Maxim Zvyagin*, Alexander Brace*, Kyle Hippe*, **Yuntian Deng***, Bin Zhang, Cindy Orozco Bohorquez, Austin Clyde, Bharat Kale, Danilo Perez-Rivera, Heng Ma, Carla M. Mann, Michael Irvin, J. Gregory Pauloski, Logan Ward, Valerie Hayot-Sasson, Murali Emani, Sam Foreman, Zhen Xie, Diangen Lin, Maulik Shukla, Weili Nie, Josh Romero, Christian Dallago, Arash Vahdat, Chaowei Xiao, Thomas Gibbs, Ian Foster, James J. Davis, Michael E. Papka, Thomas Brettin, Rick Stevens, Anima Anandkumar, Venkatram Vishwanath, Arvind Ramanathan. GenSLMs: Genome-scale language models reveal SARS-CoV-2 evolutionary dynamics. *The International Journal of High Performance Computing Applications 2023 ACM Gordon Bell Special Prize for COVID-19 Research*
- 23 **Yuntian Deng**, Noriyuki Kojima, Alexander Rush. Markup-to-Image Diffusion Models with Scheduled Sampling. *ICLR 2023 (Demo: huggingface.co/spaces/yuntian-deng/latex2im)*
- 22 Richa Rastogi, Yair Schiff, Alon Hachohen, Zhaozhi Li, Ian Lee, **Yuntian Deng**, Mert R. Sabuncu, Volodymyr Kuleshov. Semi-Parametric Inducing Point Networks and Neural Processes. *ICLR 2023*
- 21 **Yuntian Deng**, Volodymyr Kuleshov, Alexander Rush. Model Criticism for Long-Form Text Generation. *EMNLP 2022*
- 20 Anton Bakhtin*, **Yuntian Deng***, Sam Gross, Myle Ott, Marc’Aurelio Ranzato, Arthur Szlam. Residual Energy-Based Models for Text. *JMLR 2021*
- 19 Justin Chiu*, **Yuntian Deng***, Alexander Rush. Low-Rank Constraints for Fast Inference in Structured Models. *NeurIPS 2021*
- 18 **Yuntian Deng**, Alexander Rush. Sequence-to-Lattice Models for Fast Translation. *EMNLP 2021 Findings*
- 17 Keyon Vafa, **Yuntian Deng**, David Blei, Alexander Rush. Rationales for Sequential Predictions. *EMNLP 2021 Oral*
- 16 **Yuntian Deng**, Alexander Rush. Cascaded Text Generation with Markov Transformers. *NeurIPS 2020*
- 15 Thierry Tambe, En-Yu Yang, Zishen Wan, **Yuntian Deng**, Vijay Janapa Reddi, Alexander Rush, David Brooks, Gu-Yeon Wei. Algorithm-Hardware Co-Design of Adaptive Floating-Point Encodings for Resilient Deep Learning Inference. *DAC 2020 Best Paper Award*
- 14 **Yuntian Deng**, Anton Bakhtin, Myle Ott, Arthur Szlam, Marc’Aurelio Ranzato. Residual Energy-Based Models for Text Generation. *ICLR 2020*
- 13 Zachary Ziegler*, **Yuntian Deng***, Alexander Rush. Neural Linguistic Steganography. *EMNLP 2019 Oral (Demo: steganography.live)*
- 12 **Yuntian Deng**, David Rosenberg, Gideon Mann. Challenges in End-to-End Neural Scientific Table Recognition. *ICDAR 2019*
- 11 Anton Bakhtin, Sam Gross, Myle Ott, **Yuntian Deng**, Marc’Aurelio Ranzato, Arthur Szlam. Real or Fake? Learning to Discriminate Machine from Human Generated Text. *arXiv 2019*
- 10 **Yuntian Deng***, Yoon Kim*, Justin Chiu, Demi Guo, Alexander Rush. Latent Alignment and Variational Attention. *NeurIPS 2018*
- 9 Sebastian Gehrmann, **Yuntian Deng**, Alexander Rush. Bottom-Up Abstractive Summarization. *EMNLP 2018*
- 8 **Yuntian Deng**, Anssi Kanervisto, Jeffrey Ling, Alexander Rush. Image-to-Markup Generation with Coarse-to-Fine Attention. *ICML 2017 (Demo: im2markup.yuntiandeng.com)*
- 7 Pengtao Xie, **Yuntian Deng**, Yi Zhou, Abhimanu Kumar, Yaoliang Yu, James Zou, Eric P. Xing. Learning Latent Space Models with Angular Constraints. *ICML 2017*

- 6 Guillaume Klein, Yoon Kim, **Yuntian Deng**, Jean Senellart, Alexander M. Rush. OpenNMT: Open-Source Toolkit for Neural Machine Translation. *ACL 2017 Demo Best Demonstration Paper Runner-Up*
- 5 Zichao Yang, Zhiting Hu, **Yuntian Deng**, Chris Dyer, Alex Smola. Neural Machine Translation with Recurrent Attention Modeling. *EACL 2017*
- 4 Xuezhe Ma, Yingkai Gao, Zhiting Hu, Yaoliang Yu, **Yuntian Deng**, Eduard Hovy. Dropout with Expectation-linear Regularization. *ICLR 2017*
- 3 Hao Zhang, Zhiting Hu, **Yuntian Deng**, Mrinmaya Sachan, Zhicheng Yan, Eric P. Xing. Learning Concept Taxonomies from Multi-modal Data. *ACL 2016*
- 2 Pengtao Xie, **Yuntian Deng**, Eric P. Xing. Diversifying Restricted Boltzmann Machine for Document Modeling. *KDD 2015*
- 1 Zhiting Hu, Poyao Huang, **Yuntian Deng**, Yingkai Gao, Eric P. Xing. Entity Hierarchy Embedding. *ACL 2015*

Research Funding

Laude Institute Slingshots — USD 45,000 in-kind support <i>Project:</i> NeuralOS	Jun 2026
Google Research Award — co-PI; USD 25,000 unrestricted gift and USD 75,000 in Google Cloud credits <i>Project:</i> LLM-Guided Interactive Training of Biomedical Foundation Models on TPUs	Apr 2026
CIFAR Canadian AI Safety Institute – AI Safety Solution Network — co-PI; CAD 700,000 shared network funding <i>Project:</i> Safeguarding Courts from Synthetic AI Content	2025–2027
Canada Foundation for Innovation John R. Evans Leaders Fund — PI; CAD 244,846 total project value, including CAD 96,000 CFI contribution <i>Project:</i> Progressive Internalization of Skills for Advanced Language Model Reasoning	2024–2028
Digital Research Alliance of Canada, Resources for Research Groups — PI; RRG 2025 No. 5275, estimated in-kind value CAD 61,165; RRG 2026 ID 3797, estimated in-kind value CAD 28,772 <i>Project:</i> Neural OS: Simulating Operating Systems using Diffusion Models	Jul 2025–Apr 2027
Mitacs Business Strategy Internship — PI; CAD 10,000 <i>Project:</i> Decentralized Electric Microgrid Optimization by Reinforcement Learning	2025
Manulife Research Grant — PI; CAD 30,000 <i>Project:</i> End-to-End Optimization of LLM-Based Multi-Agent Systems via Text Differentiation	2024–2026
NSERC Discovery Grant — PI; RGPIN-2024-05178; CAD 41,000/year + CAD 12,500 Discovery Launch Supplement <i>Project:</i> Multi-Agent Communication and Collaboration for Specialized Large Language Models	Aug 2024–Apr 2029
University of Waterloo Starter Grant — CAD 150,000	Aug 2024–Aug 2029

Professional Service

Editorial and Leadership

Co-Chief Technical Officer, ACL Rolling Review (ARR) — technical infrastructure supporting 17,087 submissions in the May 2026 cycle	Dec 2025–Present
Action Editor, Transactions of the Association for Computational Linguistics (TACL)	2025–Present
Co-Chair, AAAI-27 Undergraduate Consortium	2027
Mentor Chair, ICLR	2026
Workshop Organizer, Latent & Implicit Thinking – Going Beyond CoT Reasoning, ICLR	Apr 27, 2026

Area Chair and Senior Program Committee Roles

Area Chair: ICLR (2024–2026), ICML (2026), NeurIPS (2026), ACL Rolling Review (2024–2026), COLM (2025–2026), EMNLP Demo (2026), LREC (2026), EMNLP GEM Workshop (2023), and IJCNLP–AAACL (Language Modeling, 2025).

Senior Program Committee: AAAI 2027.

Selected Reviewing and Selection Service

Reviewer, <i>Nature</i>	2026
Reviewer, Schmidt Sciences Science of Trustworthy AI	2026
Journals: TACL, TMLR, Computational Linguistics, ACM AI Letters, ACM Computing Surveys, Foundations and Trends in Signal Processing, IEEE TNNLS, IEEE TIFS, ACM Transactions on Interactive Intelligent Systems, ACM Transactions on Information Systems, and Journal of Computer Science and Technology.	
Conferences: ICML, NeurIPS, ICLR, ACL, ACL Rolling Review, EMNLP, AAAI, COLM, COLING, AAACL, EACL, NAACL, ACL Demo, NAACL Demo, EMNLP Demo, and various workshops and reproducibility challenges.	
Proposal Reviewer, Mitacs	
Selection Committee Member, Vector Scholarship in Artificial Intelligence	2026–2027

Conference Service

Oral Session Moderator, ICLR	2024, 2026
Oral Session Chair, EMNLP	2024
Co-host, ICLR Mentorship Session	Apr 23, 2026
Earlier conference volunteer service: ACL (2020–2021), EMNLP (2021), and ICML (2016, 2020).	

Reviewing Recognition

Top 10% Reviewer, NeurIPS	2020
Best Reviewer (Top 10%), ICML	2021

Teaching Experience

Waterloo

Instructor, CS 486/686: Introduction to Artificial Intelligence	Spring 2026
Instructor, CS 486/686: Introduction to Artificial Intelligence	Spring 2025
Guest Lecturer, Computational Linguistics (Freda Shi)	Feb 11, 2025
Instructor, CS 886: Topics in Language Models	Fall 2024

Harvard University

Head TA, Intro to Computational Linguistics and NLP (Stuart Shieber)	Fall 2020; Fall 2021
TA, Natural Language Processing (Alexander Rush)	Jan 2019–May 2019

Guest Lecturer, Natural Language Processing – Translation (Alexander Rush)	Feb 13, 2019
TA, Advanced Machine Learning (Alexander Rush)	Aug 2017–Dec 2017
Cornell University	
Course Materials Contributor, Break Through Tech AI (Alexander Rush)	Aug 2021
Carnegie Mellon University	
TA, Probabilistic Graphical Models (Matthew Gormley, Eric Xing)	Jan 2016–May 2016
TA, Graduate Machine Learning (Ziv Bar-Joseph, Eric Xing)	Aug 2015–Dec 2015
Community Service	
Co-Treasurer, Board of Directors, Bright Starts Early Learning Centre	May 2025–May 2027
Childcare Chair, NAACL 2022 Diversity & Inclusion Committee	Apr–Jul 2022
Leader, Harvard English Language Table	Sep 2021–Jan 2023
Mentor, Harvard Women in CS NLP Reading Group	Sep 2020–Sep 2023
Volunteer Instructor of AP CS, Revere High School	Sep 2020–May 2021
Research Internships	
NVIDIA (Anima Anandkumar, Weili Nie, Arash Vahdat, Chaowei Xiao)	May–Dec 2022
Facebook AI Research (Marc’Aurelio Ranzato, Arthur Szlam)	May–Dec 2019
Bloomberg CTO Office (David Rosenberg, Gideon Mann)	Jan–Aug 2017
UCSD (Charles Elkan)	Jul–Sep 2013
Talks	
Invited Talks and Seminars	
Apple MLR Internal Reading Group: NeuralOS	Aug 2026, Remote
Vector Endless Summer School: NeuralOS	Jun 2026, Remote
NSERC CREATE Responsible AI consortium: NeuralOS	Mar 2026, Remote
Mila/McGill NLP Reading Group: Implicit Reasoning in Language Models	Oct 2025, Remote
University of Cambridge Language Technology Lab: Implicit Chain-of-Thought: Internalizing Reasoning in Language Models	Oct 2025, Remote
Google DeepMind: Implicit Reasoning in Language Models	May 2025, Remote
Colin Raffel Group: Implicit Reasoning in Language Models	Mar 2025, Toronto, Canada
University of Oklahoma: Implicit Chain of Thought Reasoning	Jan 2025, Remote
NYU: Implicit Chain of Thought Reasoning	Nov 2024, New York, NY
Google DeepMind: Implicit Chain of Thought Reasoning	Jul 2024, Remote
Queen’s University: Implicit Chain of Thought Reasoning	Feb 2024, Remote
UBC: Implicit Chain of Thought Reasoning	Jan 2024, Vancouver, Canada
AI2 Mosaic: WildChat	Sep 2023, Seattle, WA
AI2 Semantic Scholar: Structure Modeling in Language Models	Aug 2023, Seattle, WA
NYU AI School 2023: Natural Language Processing	Jun 2023, New York, NY
AI2 Mosaic: Structure Modeling in Language Models	Jun 2023, Seattle, WA
UMass Amherst: Structural Coherence in Text Generation	May 2023, Remote
Georgia Tech: Structural Coherence in Text Generation	Apr 2023, Atlanta, GA

Harvard Kempner Institute: Structural Coherence in Text Generation	Mar 2023, Remote
University of Alberta: Structural Coherence in Text Generation	Mar 2023, Edmonton, Canada
TTIC: Structural Coherence in Text Generation	Feb 2023, Chicago, IL
University of Waterloo ECE: Structural Coherence in Text Generation	Feb 2023, Remote
Cornell Seminar in NLU: Structural Coherence in Text Generation	Feb 2023, New York, NY
University of Waterloo CS: Structural Coherence in Text Generation	Jan 2023, Waterloo, Canada
University of Chicago Rising Stars in Data Science: Model Criticism for Long-Form Text Generation	Nov 2022, Chicago, IL
Princeton NLP: Model Criticism for Long-Form Text Generation	Nov 2022, Princeton, NJ
OpenAI: Residual Energy-based Models for Text Generation	Apr 2021, Remote
Baidu: Cascaded Text Generation with Markov Transformers	Jun 2020, Remote
LinkedIn: Residual Energy-based Models for Text Generation	Mar 2020, Remote
Wayfair: Neural Encoder-Decoder Models	Nov 2019, Boston, MA
FAIR NLP Meeting: Energy-Based Models for Text Generation	Aug 2019, New York, NY
The French National Center for Scientific Research: Variational Attention	Jul 2018, France
Tencent Social Network Group TSAIC: OpenNMT	Dec 2017, Shenzhen, China

Conference and Workshop Presentations

EMNLP 2021 Findings: Sequence-to-Lattice Models for Fast Translation	Nov 2021, Remote
NeurIPS 2020: Cascaded Text Generation with Markov Transformers	Dec 2020, Remote
ICLR 2020: Residual Energy-based Models for Text Generation	Apr 2020, Remote
EMNLP 2019: Neural Linguistic Steganography	Nov 2019, Hong Kong
NeurIPS 2018: Latent Alignment and Variational Attention	Dec 2018, Montreal, Canada
Association for Machine Translation in the Americas: OpenNMT	Mar 2018, Boston, MA
Open-Source Neural Machine Translation Workshop: Image-to-Text	Mar 2018, Paris, France
ICML: Image-to-Markup Generation	Aug 2017, Sydney, Australia

Panels

KAIST at ICML	Jul 6, 2026, Seoul
Panelist on vision, multimodal AI, and agents.	
Vector Institute's Remarkable 2025	Mar 2025, Toronto
Panelist on emerging questions in NLP.	
ACL Mentorship: Sharing Learnings & Identifying Research Directions	Dec 2023, Singapore
Panelist on promising research directions.	
ACL Mentorship: Preparing Your Grad School Application	Oct 2023, Remote
Mentorship panelist on graduate-school applications.	
NAACL 2022: Balancing Professional and Family Commitments	Jul 2022, Remote
Co-host on balancing professional and family commitments.	

Graduate Examination and Milestone Service

All service listed below is at the University of Waterloo; students' advisers are shown in parentheses unless my supervisory role is noted.

Thesis Supervision and Examination Committees

Dongfu Jiang — PhD Thesis Committee (Wenhu Chen)	Oct 8, 2026
Henry Lin — MMath Thesis Supervision	Aug 17, 2026
Han Zhou — MMath Thesis Committee (Yang Lu)	Aug 17, 2026
Guilherme Campos — MASc Thesis Committee (Wojciech Golab and Tahsin Reza)	Aug 12, 2026
Benjamin Schneider — MMath Thesis Committee (Wenhu Chen and Florian Kerschbaum)	Aug 11, 2026
Arthur Chen — MMath Thesis Committee (Victor Zhong)	Aug 4, 2026
Ruotian Wu — MMath Thesis Committee (Pascal Poupart)	Aug 4, 2026
William Loh — PhD Thesis Committee (Pascal Poupart)	Jul 21, 2026
Xueguang Ma — PhD Thesis Committee (Jimmy Lin)	Jul 15, 2026
Hosna Oyarhoseini — MMath Thesis Committee (Jimmy Lin and Amir-Hossein Karimi)	Jul 14, 2026
Xiang Li — PhD Comprehensive-II Committee (Jian Zhao and Kate Larson)	Jun 26, 2026
Justin Xu — MMath Thesis Committee (Pascal Poupart)	May 2026
Saarang Agarwal — MMath Thesis Committee (Mei Nagappan and Pengyu Nie)	Apr 2026
Weiming Ren — PhD Comprehensive-II Committee (Wenhu Chen and Jimmy Lin)	Apr 2026
Dongfu Jiang — PhD Comprehensive-II Committee (Wenhu Chen)	Apr 3, 2026
Cong Wei — PhD Comprehensive-II Committee (Wenhu Chen)	Mar 2026
Yuansheng Ni — MMath Thesis Committee (Wenhu Chen)	Nov 2025
Xueguang Ma — PhD Comprehensive-II Committee (Jimmy Lin)	Oct 2025
Allen Zhang — MMath Thesis Committee (Anita Layton)	Aug 2025
Carter Blair — MMath Thesis Committee (Kate Larson and Edith Law)	Jun 2025
Achint Soni — MMath Thesis Committee (Sirisha Rambhatla and Charles Clarke)	May 2025
Chi-Chung Cheung — MMath Thesis Committee (Anita Layton)	Dec 2024
William Loh — PhD Thesis Proposal Committee (Pascal Poupart)	Aug 2024

PhD Seminar Service

Faculty review and approval of Waterloo's required PhD seminars; each seminar requires at least two faculty sign-offs.

Dongfu Jiang — 2 seminars (Wenhu Chen)	Jul 15 and Aug 5, 2026
Cong Wei — 2 seminars (Wenhu Chen)	Jul 24 and 27, 2026
Weiming Ren — 2 seminars (Wenhu Chen)	Jul 15 and 22, 2026
William Loh — 3 seminars (Pascal Poupart)	Mar 13; Jun 5 and 15, 2026
Ronak Pradeep — 1 seminar (Jimmy Lin)	Oct 9, 2025

Open Source Projects

ProgramAsWeights (PAW) (230 GitHub stars; 3,856 community-built neural programs)
English ↔ Claudish (two local neural functions compiled with ProgramAsWeights; 1 million+ views and 17,000+ likes on X)
NeuralOS (202 GitHub stars; 50,000+ demo visits)
WildChat-4.8M-Full (909 total Hugging Face likes across five AllenAI versions)
OpenAIWatch (101,000 visits since 2025; 87,000 in 2026 YTD)
WildVisualizer (3.1 million visits since 2025; 2.0 million in 2026 YTD)
Internalize CoT Step by Step (208 GitHub stars)
Implicit Chain of Thought (149 GitHub stars)
WildChat Chatbots (o1, o1-mini, GPT-4, and ChatGPT; 1,616 total Hugging Face likes)
Cascaded Generation (130 GitHub stars)
Neural Linguistic Steganography (14,200 demo visits since 2025; 214 GitHub stars)
Variational Attention (327 GitHub stars)
Image-to-Markup Generation (1,300+ GitHub stars)
OpenNMT-py (7,000+ GitHub stars)
OpenNMT (archived) (2,400+ GitHub stars)
Attention OCR (1,100+ GitHub stars)