

# CS 486/686

# Constraint Satisfaction Problems

Yuntian Deng

---

Lecture 4

RN 6.1–6.3 · PM 4.1–4.4

# Reminders

**DUE**

**Tue May 26**

**Chat 1** (Intro & uninformed search) — due 11:59 PM.

**RELEASED TODAY**

**Due Thu May 28**

**Chat 2** (Heuristic search & CSPs) — two-day extension from the original Tue May 26.

**DISTINGUISHED LECTURE**

**Tue May 26, 10:00–11:30 AM · DC 1302**

**Prof. [Kyunghyun Cho](#)** (NYU) — co-inventor of *attention* (Bahdanau, Cho, Bengio, ICLR 2015) and the *GRU*. **Highly recommended.**

Search ›

Uncertainty ›

Decisions ›

Learning

## Learning goals

- Formulate a real-world problem as a CSP
- Trace backtracking search
- Verify whether a constraint is arc-consistent
- Trace the AC-3 algorithm
- Trace backtracking *combined* with arc consistency


# CSPs you've seen

|   |   |   |   |   |
|---|---|---|---|---|
| C | A | T |   | S |
|   |   |   |   |   |
| A | R | E |   |   |
|   |   |   |   |   |
| D |   | O | N | E |


Crosswords

|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 5 | 3 |   | 7 |   |   |   |
| 6 |   |   | 1 | 9 | 5 |   |
|   | 9 | 8 |   |   |   | 6 |
| 8 |   |   | 6 |   |   | 3 |
| 4 |   | 8 |   |   | 3 |   |
| 7 |   |   | 2 |   |   | 6 |
|   | 6 |   |   |   |   | 2 |
|   |   |   | 4 | 1 | 9 |   |
|   |   |   |   | 8 |   |   |
|   |   |   |   |   | 7 | 9 |

Sudoku



Graph colouring



N-queens

# Naive idea: generate-and-test

Enumerate every assignment, test if it satisfies the constraints.

```
A=1, B=1, C=1, D=1, E=1    # fail
```

```
A=1, B=1, C=1, D=1, E=2    # fail
```


```
...
```

5 variables, domain size 4:  $4^5 = 1024$  assignments. Doesn't scale.

And it's wasteful: many constraints can be checked on *partial* assignments.

# Search algorithms are structure-blind

A search algorithm sees only "generate successors" and "is this a goal?".



Two queens already in the same row — clearly a dead end — yet the search keeps generating successors. We need to **see the structure**.

A CSP is a triple  $(X, D, C)$

**Variables  $X$**

$\{X_1, X_2, \dots, X_n\}$

**Domains  $D$**

Each  $X_i$  draws values  
from a finite set  $D_i$ .

**Constraints  $C$**

Each  $c \in C$  restricts  
allowed combinations  
of variable values.

**Solution:** an assignment to all variables that satisfies every constraint.

# 4-queens: variables and domains

- One queen per column  $\Rightarrow$  track its *row*.
- **Variables:**  $x_0, x_1, x_2, x_3$  (column index).
- **Domains:**  $D_{x_i} = \{0, 1, 2, 3\}$ .

|   |       |       |       |       |
|---|-------|-------|-------|-------|
|   | $x_0$ | $x_1$ | $x_2$ | $x_3$ |
| 0 |       |       |       |       |
| 1 |       |       |       |       |
| 2 |       |       |       |       |
| 3 |       |       |       |       |



# 4-queens: constraints

No pair of queens in the same row or diagonal.

$$\forall i \neq j : (x_i \neq x_j) \wedge (|x_i - x_j| \neq |i - j|)$$

$$\text{Columns } 0, 1: (x_0 \neq x_1) \wedge (|x_0 - x_1| \neq 1).$$

No *column* constraint needed — our variables already encode “one queen per column”.

# Q: which constraints do we need?

Q. Given the variables  $x_0, \dots, x_3$  and their domains, which row/column/diagonal constraints must we enforce?

- A. No two queens in the same row
- B. No two queens in the same column
- C. No two queens on the same diagonal
- D. Two of (A), (B), (C)
- E. All of (A), (B), (C)

**D — row + diagonal only.** The column constraint is already implied by our choice of variables: there's exactly one  $x_i$  per column, so no two queens can ever share a column.

# Two ways to express a constraint

Constraint: queens in columns 0 and 2 are not in the same row or diagonal.

## As a formula

$$(x_0 \neq x_2) \wedge (|x_0 - x_2| \neq 2)$$

## As a table

| $x_0$ | $x_2$ |
|-------|-------|
| 0     | 1     |
| 0     | 3     |
| 1     | 0     |
| 1     | 2     |
| 2     | 1     |
| 2     | 3     |
| 3     | 0     |
| 3     | 2     |

## Q: encode $(x_0, x_2)$ as a formula

Q. Which formula encodes "queens in columns 0 and 2 are not in the same row or diagonal"?

- A.  $(x_0 \neq x_2)$
- B.  $(x_0 \neq x_2) \wedge ((x_0 - x_2) \neq 1)$
- C.  $(x_0 \neq x_2) \wedge ((x_0 - x_2) \neq 2)$
- D.  $(x_0 \neq x_2) \wedge (|x_0 - x_2| \neq 1)$
- E.  $(x_0 \neq x_2) \wedge (|x_0 - x_2| \neq 2)$

**E.** The two columns are 2 apart, so diagonal conflict means  $|x_0 - x_2| = 2$ . We need both the row and diagonal exclusions, and the absolute value because the queens could be above or below each other.

# 4-queens as an incremental search problem

- **State:** one queen per column in the leftmost  $k$  columns, no pair attacking each other.
- **Initial state:** empty board \_ \_ \_ \_.
- **Goal state:** 4 queens placed, no pair attacking. e.g. 1 3 0 2.
- **Successor:** place a queen in the leftmost empty column that no current queen attacks.

e.g. 0 \_ \_ \_  $\rightarrow$  0 2 \_ \_ , 0 3 \_ \_ .

Any search algorithm works on this state space. **Backtracking** = **DFS** + one rule: skip any partial assignment that already violates a constraint.

# Backtracking search

**Backtracking = DFS** through partial assignments + prune any partial that already violates a constraint.


## **backtrack(assignment, csp)**

1. If **assignment** is complete → **return** it.
2. Pick an unassigned variable **var**.
3. For each **value** in **domain(var)** that satisfies every constraint:
4. Add **{var = value}** and go deeper; if a solution is found, return it.
5. Otherwise undo and try the next value.
6. If no value worked → **return failure**.

# Trace: backtracking on 4-queens

Each node = partial assignment. Yellow = dead end. Green = solution.

**Step:**



Found in 9 expansions.

# Partial assignment, fatal future

After placing  $x_0 = 0$ , is  $x_1 = 2$  safe?

|   |   |   |   |   |
|---|---|---|---|---|
| 0 | Q | X | X | X |
| 1 | X | X | X |   |
| 2 | X | Q | X | X |
| 3 | X | X | X | X |

Every cell of column  $x_3$  is attacked  $\rightarrow$  **no future** for the remaining queens.

Knowing this *before* trying  $x_2 =$  **arc consistency**.




## Binary constraints only

Arc consistency is defined on *binary* constraints  $c(X, Y)$ .

**Unary** constraints  $c(X)$ : drop disallowed values directly from  $D_X$ .

The CSPs in this course (4-queens, the homework problems) are all binary.

# Notation: arcs



- An **arc**  $\langle X, c(X, Y) \rangle$  pairs a constraint with one of its variables.
- $X$  is the **primary** variable (the one we'll prune);  $Y$  is the **secondary**.
- Each binary constraint defines *two* arcs — one per primary choice.

# Arc-consistency definition

**Arc consistency.** The arc  $\langle X, c(X, Y) \rangle$  is *arc-consistent* iff for every  $v \in D_X$  there exists a  $w \in D_Y$  such that  $(v, w)$  satisfies  $c(X, Y)$ .

**Intuition:** every value in the primary domain has a partner in the secondary domain.

If some  $v \in D_X$  has no partner, remove  $v$  from  $D_X$  to make the arc consistent.

## Q: is this arc consistent?

Q. Constraint  $X < Y$ , with  $D_X = D_Y = \{1, 2\}$ . Is  $\langle X, X < Y \rangle$  arc-consistent?

A. Yes

B. No

**B — No.** For  $v = 2 \in D_X$ , there is no  $w \in D_Y$  with  $2 < w$ . We can **remove 2 from  $D_X$**  to make the arc consistent.

# The AC-3 algorithm


Process arcs one at a time. After each shrink, re-queue neighbours whose consistency may have broken.

## AC-3(arcs)

1.  $S \leftarrow$  set of all arcs. *initialise the queue*
2. While  $S$  is not empty:
  1. Pop an arc  $\langle X, c(X, Y) \rangle$  from  $S$ .
  2. For each  $v \in D_X$  with no  $w \in D_Y$  satisfying  $c(X, Y) \rightarrow$  remove  $v$  from  $D_X$ .
  3. If  $D_X$  is now empty  $\rightarrow$  return **false**. *infeasible — no value works*
  4. If  $D_X$  shrank  $\rightarrow$  for every neighbour  $Z \neq Y$ , add  $\langle Z, c'(Z, X) \rangle$  back into  $S$ . *re-queue — its consistency may have broken*
3. Return **true**. *every arc is consistent*

# Why re-queue arcs?

Shrinking one domain can break a previously-consistent neighbouring arc.




Initially: arc  $\langle X, X < Z \rangle$  is consistent (1<2, 1<3, 2<3 all work).

**Then:** arc  $\langle X, Y = X \rangle$  forces  $D_X \rightarrow \{2\}$ .

**Now:**  $\langle Z, X < Z \rangle$  is broken — for  $w = 2 \in D_Z$ , no  $v < 2$  remains in  $D_X$ . **Must re-queue.**

# Trace: AC-3 on 4-queens with $x_0 = 0$



## Three outcomes of AC-3

- A domain becomes empty → no solution.
- Every domain has exactly one value → solution found, no search needed.
- Some domains have multiple values → need backtracking search on what remains.




# AC-3 — properties

|                 |                                                                                                                           |
|-----------------|---------------------------------------------------------------------------------------------------------------------------|
| Terminates?     | <b>Yes</b> <i>each arc enqueued at most <math>d</math> times</i>                                                          |
| Order of arcs?  | Doesn't matter <i>same final domains</i>                                                                                  |
| Time complexity | $O(c d^3)$ <i><math>c</math> constraints <math>\times d</math> re-queues per arc <math>\times O(d^2)</math> per check</i> |

# Combining backtracking with AC-3

1. Run backtracking search.
2. After each value assignment, run AC-3.
3. If a domain becomes empty → **fail this branch**, backtrack.
4. If every domain has exactly one value → **solution found**.
5. Otherwise → continue backtracking on the remaining unassigned variables.

# Trace: backtracking + AC-3 on 4-queens



## Why combine BT with AC-3?

- **AC-3 prunes branches BT would otherwise explore** — it kills bad partial assignments before they grow.
- **BT decides what AC-3 can't** — when AC-3 leaves multi-value domains, BT picks a value and triggers another AC-3 pass.
- On 4-queens with  $x_0 = 1$ , AC-3 alone reduced everything to a single solution — **zero further search**.

# Three techniques for CSPs

| Technique    | How                                                    | Finds a solution?                            | Best at                        |
|--------------|--------------------------------------------------------|----------------------------------------------|--------------------------------|
| Backtracking | DFS through assignments, prune on constraint violation | Yes                                          | generic, easy to implement     |
| AC-3 alone   | Repeatedly enforce arc consistency                     | <i>only if all domains become singletons</i> | shrinking domains preemptively |
| BT + AC-3    | BT picks a value, AC-3 propagates                      | Yes <i>far fewer expansions</i>              | practical CSP solving          |

## Learning goals (recap)

- ✓ Formulate a real-world problem as a CSP
- ✓ Trace backtracking search
- ✓ Verify arc consistency
- ✓ Trace AC-3
- ✓ Combine backtracking with arc consistency