

# CS 486/686

# Local Search

Yuntian Deng

---

Lecture 5

RN 4.1 · PM 4.7–4.8

# Reminders

**DUE TONIGHT**

Tue May 26, 11:59 PM

**Chat 1** (Intro & uninformed search) on Chrysalis.

**DUE THU**

Thu May 28, 11:59 PM

**Chat 2** (Heuristic search & CSPs) — extended from Tue May 26.

Search ›

Uncertainty ›

Decisions ›

Learning

## Learning goals

- Describe the advantages of local search
- Formulate a real-world problem as local search
- Verify whether a state is a local / global optimum
- Describe strategies for escaping local optima
- Trace greedy descent, random restarts, simulated annealing, genetic algorithms
- Compare and contrast local-search algorithms

# Why local search?

Classical search algorithms...

- ... explore the space **systematically**. *What if the space is too large, or infinite?*
- ... remember a **path** from the initial state. *What if we only care about the final state (e.g. CSP solution)?*

**Solution: local search.**

## Local search: what it gives up

- No systematic exploration of the space
- No path back to the start
- No proof when no solution exists

## Local search: what it buys you

- Fast, good-enough states on average
- Very small memory footprint
- Handles pure optimization (no goal test needed)

# What is local search?

Start with a **complete** assignment to all variables. Take iterative steps to improve it.

## State

A complete assignment to all variables.

## Neighbour relation


Which states can we step to next?

## Cost function

How good is each state?  
Lower is better.

# 4-queens as local search

- **State:** one queen per column, any row — e.g. (3, 2, 1, 1).
- **Initial:** random state.
- **Goal:** no attacking pair.
- **Neighbours:** (A) move one queen in its column; (B) swap two queens' rows.
- **Cost:** number of attacking pairs.



|   |   |   |   |   |
|---|---|---|---|---|
| 0 |   |   |   |   |
| 1 |   |   | Q | Q |
| 2 |   | Q |   |   |
| 3 | Q |   |   |   |


(3, 2, 1, 1), cost = 4.

# Greedy descent

a.k.a. hill climbing

- **Start** from a random state.
- **Move** to the neighbour with the lowest cost — *if it improves on the current state*.
- **Stop** when no neighbour is better.

# Greedy descent on a landscape



## Greedy descent: in a sentence


*Descend into a canyon in a thick fog.*

- **Descend:** always move to the best neighbour.
- **Thick fog:** only see immediate neighbours.

Fast in practice — but only finds **local** optima.

# State-space landscape

**Local optimum:** no neighbour has strictly lower cost. **Global optimum:** lowest cost overall.



---

A flat region with no slope is a *plateau*; a flat region at the top of a slope is a *shoulder*.

# Q: local or global optimum?

**Q.** Consider neighbour relation **B**: swap the row positions of two queens. Is this state a local / global optimum?

|   |   |   |   |
|---|---|---|---|
|   |   | Q |   |
|   |   |   | Q |
|   | Q |   |   |
| Q |   |   |   |


State  $(x_0, x_1, x_2, x_3) = (3, 2, 0, 1)$ . Cost = 2.

- A. Local optimum and global optimum.
- B. Local optimum, but not global.
- C. Not a local optimum, but it is the global optimum.
- D. Neither a local nor a global optimum.

**D — neither.** Swap  $x_0$  and  $x_2$ : new state  $(0, 2, 3, 1)$  has cost 1 (better). So not a local

# Q: local or global optimum?

**Q.** Same state, but now neighbour relation **A**: move *one* queen to another row in the same column.



State  $(3, 2, 0, 1)$ . Cost = 2.

- A. Local optimum and global optimum.
- B. Local optimum, but not global.
- C. Not a local optimum, but it is the global optimum.
- D. Neither a local nor a global optimum.

**B — local optimum, but not global.** Every single-queen move keeps or worsens the cost (try

# Escaping flat local optima

## Sideway moves

Accept neighbours of *equal* cost. Cap consecutive sideways to avoid loops.

## Tabu list

Forbid recently-visited states — short-term memory.

# Sideway moves help — a lot

8-queens:  $\approx 17$  million states.

| Algorithm                         | % solved | Steps to success | Steps to failure |
|-----------------------------------|----------|------------------|------------------|
| Greedy descent                    | 14%      | 3–4              | 3–4              |
| Greedy + $\leq 100$ sideway moves | 94%      | 21               | 64               |

A tiny tolerance for flat moves turns 14% into 94%.

# Choosing the neighbour relation

Aim for a **small incremental change** to the variable assignment.

## Bigger neighbourhoods

- Compare more nodes per step
- More likely to find the best step
- Each step takes more time

## Smaller neighbourhoods

- Compare fewer nodes per step
- May miss the best step
- Each step is fast

# Greedy descent gets stuck. What now?

## Random restarts

Jump to a fresh random state.

*e.g. greedy descent with restarts.*

## Random walks

Sometimes step to a *random* neighbour, not just a better one.

*e.g. simulated annealing.*

# Q: which random move helps where?



(a) mostly flat, hidden deep valleys



(b) bumpy bowl with one deep valley

Q. Which random move is better for each landscape?

**(a) Random restarts** — basins are widely separated; a walk can't drift between them, but a fresh restart can land in a different valley.

**(b) Random walks** — the bumps are small; a few random steps slip out of a shallow local optimum and into the global basin.

# Greedy descent with random restarts

*If at first you don't succeed, try, try again.*

- Run greedy descent multiple times, each from a fresh random state.
- Keep the best state seen across runs.
- **With enough restarts**, finds the global optimum with probability  $\rightarrow 1$  — eventually a restart will start in the global basin.

# Simulated annealing

*Annealing*: cool a molten metal slowly to make it strong.


- Start hot — high temperature  $T$ ; cool over time.
- Each step: try a *random* neighbour.
- Better? Always accept.
- Worse? Accept with probability based on  $T$  and how much worse.

High  $T$ : exploration. Low  $T$ : exploitation.

# Probability of accepting a worse neighbour

Current state  $A$ , worse neighbour  $A'$ ,  $\Delta C = \text{cost}(A') - \text{cost}(A) > 0$ ,  
temperature  $T$ :

$$p = e^{-\Delta C/T}$$



Hot temperatures stay willing to accept bad moves; cold temperatures barely budge.

## Q: as $T$ decreases...

Q. As  $T$  decreases, the probability of moving to a worse neighbour...

- A. increases
- B. stays the same
- C. decreases

**C — decreases.** Example with  $\Delta C = 10$ :  $T = 100 \Rightarrow p = e^{-0.1} \approx 0.90$ ;  $T = 10 \Rightarrow p = e^{-1} \approx 0.37$ .

## Q: as $\Delta C$ increases...

Q. As  $\Delta C$  increases (the neighbour is much worse), the probability...

- A. increases
- B. stays the same
- C. decreases

**C — also decreases.** Example with  $T = 10$ :  $\Delta C = 10 \Rightarrow p \approx 0.37$ ;  $\Delta C = 100 \Rightarrow p \approx 0.000045$ .

# Simulated annealing algorithm

## `simulated-annealing()`

1. `current`  $\leftarrow$  initial state; `T`  $\leftarrow$  large positive value.
2. While `T`  $>$  0 :
3. Sample a random neighbour `next` .
4. Let  $\Delta C = \text{cost}(\text{next}) - \text{cost}(\text{current})$ .
5. If  $\Delta C < 0 \rightarrow$  accept ( `current`  $\leftarrow$  `next` ); else accept with probability  $p = e^{-\Delta C/T}$ .
6. Cool  $T$  (e.g.  $T \leftarrow \alpha T$ ).
7. Return `current` .


# Annealing schedule

How fast should  $T$  cool?


- **In theory:** very slowly. If  $T$  decreases slowly enough, simulated annealing finds the global optimum with probability  $\rightarrow 1$ .
- **In practice: geometric cooling** —  $T_{k+1} = \alpha T_k$  with  $\alpha$  close to 1.

Example: start at  $T_0 = 10$ ,  $\alpha = 0.99$ . After 500 steps,  $T \approx 0.07$ .

# SA in action: travelling salesman



high  $T$ : tangled



low  $T$ : smooth

30 cities  $\rightarrow 30! \approx 2.6 \times 10^{32}$  tours. NP-hard — yet SA finds great tours fast.

## SA as life


- High  $T$  — young, exploring.
- Low  $T$  — older, refining.
- Exploration vs. exploitation, all over again.

## From a single state to a population

- So far, every algorithm remembers *one* state at a time.
- What if we keep  $k$  **states** in parallel?

# Beam search

- Keep the  $k$  best states in parallel.
- Expand all their neighbours.
- Keep the new  $k$  best from the pool.
- $k$  controls memory + parallelism.



Green = the  $k = 3$  kept per level.

# Beam search — corner cases

$k = 1$

greedy descent

---

$k = \infty$

breadth-first search

---

**vs.  $k$   
parallel  
restarts**

beam *shares* across threads; restarts  
are independent

---

**Watch out**

states can clump — lack of diversity

# Stochastic beam search

- Pick the next  $k$  states **probabilistically**, with probability proportional to fitness (e.g.  $\propto e^{-\text{cost}/T}$ ).
- Maintains **diversity** in the population.
- Mimics natural selection — *asexual* reproduction (mutate, keep best).

Adding *sexual* reproduction gives the **genetic algorithm**.

# Genetic algorithm

Population evolves; fittest survive; offspring combine and mutate.

## **genetic-algorithm()**

1. Initialise a population  $P$  of random individuals.
2. Until converged:
3. Select parents from  $P$  weighted by fitness.
4. Crossover parents  $\rightarrow$  offspring.
5. Mutate offspring with small probability.
6. Replace  $P$  with the offspring.
7. Return the best individual in  $P$ .

# Crossover

parent A 

parent B 

child 1 

child 2 

Split parents at a point; swap tails.

# Mutation

before **1****0****1****1****0****1****0****1**



after **1****0****1****1****1****1****0****1**

Flip a random bit with small probability.

# Four local-search algorithms

| Algorithm            | Explores how                                    | Remembers             | Finds global?                              | Best at                      |
|----------------------|-------------------------------------------------|-----------------------|--------------------------------------------|------------------------------|
| Greedy descent       | best neighbour                                  | 1 state               | <b>No</b><br><i>stuck at local optimum</i> | fast first cut               |
| GD + random restarts | best, then restart                              | 1 state + best so far | <b>Yes</b><br><i>in the limit</i>          | landscapes with wide basins  |
| Simulated annealing  | random; accept worse with $p = e^{-\Delta C/T}$ | 1 state + T           | <b>Yes</b><br><i>slow enough cooling</i>   | bumpy landscapes             |
| Genetic algorithm    | crossover + mutation                            | population of $k$     | <i>probabilistically</i>                   | complex combinatorial spaces |

# Learning goals (recap)

- ✓ Describe the advantages of local search
- ✓ Formulate a real-world problem as local search
- ✓ Verify whether a state is a local / global optimum
- ✓ Describe strategies for escaping local optima
- ✓ Trace greedy descent, random restarts, simulated annealing, genetic algorithms
- ✓ Compare and contrast local-search algorithms

## Next module: Reasoning under Uncertainty

We've optimized in the *deterministic* world. The agent will now have to reason when it can't be sure of the state or the outcome.

- L6 — probability rules: sum, product, chain, Bayes.
- L7–L8 — independence + Bayesian networks + d-separation.
- L9 — variable elimination for inference.
- L10–L11 — hidden Markov models, filtering, smoothing, Viterbi.