

CS 486/686

Variable Elimination Algorithm

Yuntian Deng

Lecture 9

RN 13.4 · PM 8.4

Reminders

 **Chrysalis is down now** (server issue), and is also scheduled offline **Thu Jun 11 (9:30 AM) – Fri Jun 12 (noon)**.

 **Deadlines extended: Chat 4 → Tue Jun 16, Chat 5 → Thu Jun 18.**

 **Assignment 1 questions → Dake Zhang on Piazza.**

Search ›

Uncertainty ›


Decisions ›

Learning

Learning goals

- Do inference in Bayes nets **more efficiently** than the joint
- Define **factors** and the four operations on them
- Trace the **variable elimination algorithm** for prior / posterior queries
- Explain how the **elimination ordering** affects complexity
- Know the basics of **approximate** inference (sampling)

A query in the Holmes BN: $P(B|w \wedge g)$



"What's the probability of a burglary, given that both Watson and Gibbon called?"

- Query variable: B
- Evidence variables: $W = w, G = g$
- Hidden variables: E, A, R (to be summed out)

Convention: capital letters are random variables, lowercase letters are specific values.

Naive: brute-force sum-out

Marginalize the full joint using the BN factorization:

$$P(B|w \wedge g) \propto \sum_e \sum_a \sum_r P(B) P(e) P(a|B \wedge e) P(w|a) P(g|a) P(r|e)$$

Q1. How many additions + multiplications does this require?

- A. ≤ 10
- B. 11–20
- C. 21–40
- D. 41–60
- E. ≥ 61

D — 47 ops. Eight terms in the product, evaluated for every $(b, e, a, r) \in \{T, F\}^4 = 16$ combinations. Lots of redundant work.

Push the sums to the right

$P(B)$ doesn't depend on e, a, r — pull it out. $P(r|e)$ only depends on e, r — push $\sum_r$ inward. And so on:

$$\begin{aligned} P(B) \sum_e P(e) \underbrace{\sum_r P(r|e)}_{=1} \sum_a P(a|B \wedge e) P(w|a) P(g|a) \\ = P(B) \sum_e P(e) \sum_a P(a|B \wedge e) P(w|a) P(g|a) \end{aligned}$$

Q2. Ops now?

B — 14 ops. Inner sum: 4 mul + 1 add; times 2 for $\sum_a$ over $\{T, F\}$; plus $\sum_e$ (4 ops); plus $P(B)$ multiplication. **~3x speedup just from re-ordering.**

The big idea

- **Reuse** intermediate computations — dynamic programming for probabilities.
- **Exploit** conditional independence baked into the Bayes net.
- Push sums in *as far as possible*, then operate on small tables of probabilities — **factors**.

This is the **Variable Elimination Algorithm**.

Factors

A **factor** $f(X_1, \dots, X_j)$ is a function from random-variable assignments to a number. It can represent a joint, a conditional, or a partial assignment.

For the Holmes BN, define one factor per CPT:

$$\begin{aligned} f_1(B) &= P(B), & f_2(E) &= P(E) \\ f_3(A, B, E) &= P(A|B \wedge E), & f_4(R, E) &= P(R|E) \\ f_5(W, A) &= P(W|A), & f_6(G, A) &= P(G|A) \end{aligned}$$

Restrict: assign a variable

Restricting $f(X_1, \dots, X_j)$ to $X_1 = v_1$ fixes that variable, leaving a smaller factor on $X_2, \dots, X_j$.

$$\text{Example: } f_1(X, Y) \rightarrow f_2(Y) = f_1(X = t, Y)$$

$f_1(X, Y)$

X	Y	val
t	t	0.1
t	f	0.9
f	t	0.2
f	f	0.8

$X=t$
 $\longrightarrow$

$f_2(Y)$

Y	val
t	0.1
f	0.9

Sum out: marginalize a variable

Summing out X_1 from $f(X_1, \dots, X_j)$ gives a factor on $X_2, \dots, X_j$ defined by

$$\left(\sum_{X_1} f\right)(X_2, \dots, X_j) = \sum_v f(X_1 = v, X_2, \dots, X_j).$$

Example: $\sum_Y f_1(X, Y) \rightarrow f_2(X)$

$f_1(X, Y)$

X	Y	val
t	t	0.1
t	f	0.4
f	t	0.3
f	f	0.5

$\xrightarrow{\sum_Y}$

$f_2(X) = \sum_Y f_1$

X	val
t	0.1 + 0.4 = 0.5
f	0.3 + 0.5 = 0.8

Multiply two factors

$(f_1 \times f_2)(X, Y, Z) = f_1(X, Y) \cdot f_2(Y, Z)$ for shared Y — the product is over the *union* of variables.

$f_1(X, Y)$			$f_2(Y, Z)$			$f_1 \times f_2$			
X	Y	val	Y	Z	val	X	Y	Z	val
t	t	0.1	t	t	0.3	t	t	t	0.03
t	f	0.9	t	f	0.7	t	t	f	0.07
f	t	0.2	f	t	0.6	t	f	t	0.54
f	f	0.8	f	f	0.4	t	f	f	0.36
						f	t	t	0.06
						f	t	f	0.14
						f	f	t	0.48
						f	f	f	0.32

Normalize a factor

Divide every entry by the sum, so the values sum to 1 (a valid probability distribution).

$f(Y)$

Y	val
t	0.2
f	0.6

normalize
→

$P(Y)$

Y	val
t	0.25
f	0.75

$$0.2 / (0.2 + 0.6) = 0.25, 0.6 / 0.8 = 0.75.$$

The variable elimination algorithm


To compute $P(X_q | X_{o_1} = v_1 \wedge \dots \wedge X_{o_j} = v_j)$:

variable-elimination(BN, query, evidence)

1. Construct a factor for each CPT in the BN.
2. Restrict every observed variable to its value.
3. For each hidden variable X_h (in chosen order): multiply all factors that contain X_h , then sum out X_h from the product.
4. Multiply the remaining factors.
5. Normalize the result.

Each iteration removes one hidden variable while keeping all factors small.

Worked example: $P(B|\neg a)$



Sub-BN for the example. B = query, A = false = evidence.

Given CPTs:

$$P(B) = 0.3, \quad P(E) = 0.1$$

$$P(A|E, B) = 0.8, \quad P(A|E, \neg B) = 0.2$$

$$P(A|\neg E, B) = 0.7, \quad P(A|\neg E, \neg B) = 0.1$$

$$P(W|A) = 0.8, \quad P(W|\neg A) = 0.4$$

Define factors:

$$f_1(B), \quad f_2(E), \quad f_3(A, B, E), \quad f_4(W, A)$$

Eliminate hidden variables

$$f_5(B, E) = f_3(\neg a, B, E)$$

B	E	val
t	t	0.2
t	f	0.3
f	t	0.8
f	f	0.9

$$f_8(B, E) = f_2 \times f_5$$

B	E	val
t	t	$0.2 \times 0.1 = 0.02$
t	f	$0.3 \times 0.9 = 0.27$
f	t	$0.8 \times 0.1 = 0.08$
f	f	$0.9 \times 0.9 = 0.81$

$$f_9(B) = \sum_E f_8$$

B	val
t	$0.02 + 0.27 = 0.29$
f	$0.08 + 0.81 = 0.89$

Multiply remaining factors, normalize

Remaining factors: $f_1(B)$, $f_9(B)$, $f_7() = 1.0$. Multiply them:

$$f_{10}(B) = f_1 \times f_9 \times f_7$$

B	val
t	$0.3 \times 0.29 \times 1 = 0.087$
f	$0.7 \times 0.89 \times 1 = 0.623$

normalize




$$P(B|\neg a)$$

B	val
t	$0.087 / 0.710 \approx 0.1225$
f	$0.623 / 0.710 \approx 0.8775$

Burglary is unlikely ($\approx 12\%$) given that the alarm did NOT sound.

Ordering matters



Compute $P(G)$ (no evidence). Initial factors:

$$f_1(E), f_2(E, R), f_3(B) \\ f_4(E, B, A), f_5(W, A), f_6(G, A)$$

- **Every** ordering of $\{R, W, E, B, A\}$ yields a correct algorithm.
- Different orderings produce **different intermediate factors**.
- Larger intermediate factors $\Rightarrow$ more space + ops.

Two orderings, very different costs

Good: $R \rightarrow W \rightarrow E \rightarrow$
 $B \rightarrow A$

- $f_2(E, R) \rightarrow f_7(E)$: 2 ops
- $f_5(W, A) \rightarrow f_8(A)$: 2 ops
- $f_1 \cdot f_7 \cdot f_4 \rightarrow f_9(B, A)$: 20 ops
- $f_9 \cdot f_8 \rightarrow f_{10}(A)$: 6 ops
- $f_{10} \cdot f_6 \rightarrow f_{11}(G)$: 6 ops

Total: 36 ops

Bad: $A \rightarrow B \rightarrow E \rightarrow R \rightarrow$
 W


- $f_4 \cdot f_5 \cdot f_6 \rightarrow f_7(E, B, W, G)$: 80 ops
- $f_7 \cdot f_3 \rightarrow f_8(E, W, G)$: 24 ops
- $f_1 \cdot f_8 \cdot f_2 \rightarrow f_9(W, G, R)$: 24 ops
- $f_9 \rightarrow f_{10}(W, G)$: 4 ops
- $f_{10} \rightarrow f_{11}(G)$: 2 ops

Total: 134 ops

The bad ordering builds a 4-variable intermediate factor $f_7(E, B, W, G)$ — that's what we pay for.

Elimination width + tree width

- **Elimination width:** the largest factor (hyperedge) an ordering ever builds.
- **Tree width** ω = smallest width over all orderings; VE costs $2^{O(\omega)}$ time and space.
- Optimal ordering is **NP-hard**, but heuristics do well ($\omega \approx 8\text{--}10$ even with 1000s of variables).




Hyperedges (colored bubbles) = factors.

Polytrees: an easy case


A **polytree** is a singly-connected BN: at most one undirected path between any two nodes. (Nodes may still have multiple parents.)

- Heuristic: **always eliminate a singly-connected node.**
- A singly-connected node always exists $\Rightarrow$ factor size *never grows*.
- Tree width = max # of parents of any node.
- VE runs in time **linear in the size of the network**.



Eliminate leaves D, E, F first — factors never grow beyond original CPTs.

Forward sampling: estimate $P(G)$




When the BN is too large for exact VE, just **simulate** it.

```
for i in 1..n: ei ~ P(E) bi ~  
P(B) ai ~ P(A | E=ei, B=bi) gi  
~ P(G | A=ai) P(G = T) ≈ (1/n)  
Σi gi
```

Sample in **topological order** so every CPT input is already known. Estimate converges to true $P(G)$ as $n \rightarrow \infty$.

Rejection sampling: estimate $P(G|\neg b)$



```
for i in 1..n: forward-sample  
  ei, bi, ai, gi if bi == T:  
  reject sample else: keep  $\tilde{g}_i$   
 $P(G = T \mid \neg b) \approx (\sum_i \tilde{g}_i) / \tilde{n}$ 
```

- Only "accepted" samples (consistent with evidence) count.
- **Wasteful** when evidence is rare — most samples thrown away.

Learning goals (recap)

- ✓ Do inference in Bayes nets **more efficiently** than the joint
- ✓ Define **factors** and the four operations on them
- ✓ Trace the **variable elimination algorithm**
- ✓ Explain how the **elimination ordering** affects complexity
- ✓ Know the basics of **approximate** inference

Next: Hidden Markov Models

A Bayes net **unrolled over time** — inference about what's happening *now* given a sequence of observations.

VEA in disguise: forward-backward, Viterbi, smoothing.