

CS 486/686

Hidden Markov Models II

Yuntian Deng

Lecture 11

RN 14.2.2

Reminders



Chat 4 is due TODAY (Tue Jun 16); Chat 5 due Thu Jun 18 — both 11:59 PM on Chrysalis.



Assignment 1 deadline extended to Thu Jun 25 (was Jun 18). The PDF, starter code, and model file were **updated on LEARN** — please re-download.



Assignment 1 questions → Dake Zhang on Piazza.

Search ›

Uncertainty ›

Decisions ›

Learning


Learning goals

- Compute the **smoothing** probability for a time step in an HMM.
- **Justify** each step in the derivation of the smoothing formula.
- Describe the **forward-backward** algorithm.
- Describe the **Viterbi** algorithm.

Recap: filtering (from L10)

From past evidence $o_{0:k}$, filtering gives the posterior over the *current* state:

$$f_{0:k} = P(S_k | o_{0:k}) = \alpha P(o_k | S_k) \sum_{s_{k-1}} P(S_k | s_{k-1}) f_{0:(k-1)}$$




- Yellow = query state S_k (the *present*).
- Gray = observed evidence.
- Cost: $O(k)$ ops via the recursive predict-update.

Today: a query about a **past** state, given all the evidence we have.

Smoothing: query a past state

$$P(S_k | o_{0:(t-1)}), \quad 0 \leq k \leq t - 1$$



Filtering (L10)

$P(S_k | o_{0:k})$ — the present state from past evidence.

Smoothing (today)

$P(S_k | o_{0:(t-1)}), k < t - 1$ — a past state from all evidence (past *and* future).

Deriving the smoothing formula

Four small steps from $P(S_k | o_{0:(t-1)})$ to the forward-backward product:

$$P(S_k | o_{0:(t-1)})$$

$$= P(S_k | o_{(k+1):(t-1)} \wedge o_{0:k})$$

rewrite

$$= \alpha P(S_k | o_{0:k}) P(o_{(k+1):(t-1)} | S_k \wedge o_{0:k})$$

Bayes' rule

$$= \alpha P(S_k | o_{0:k}) P(o_{(k+1):(t-1)} | S_k)$$

Markov

$$= \alpha f_{0:k} b_{(k+1):(t-1)}$$


definitions

Try to identify the rule before the label appears.

Smoothing = forward × backward

Split the evidence around S_k into past and future:

$$P(S_k | o_{0:(t-1)}) = \alpha \underbrace{P(S_k | o_{0:k})}_{\text{forward } f_{0:k}} \underbrace{P(o_{(k+1):(t-1)} | S_k)_{\text{backward } b_{(k+1):(t-1)}}$$



Two sweeps meet at S_k ; multiply pointwise and normalize.

Backward recursion

Compute $b_{(k+1):(t-1)} \triangleq P(o_{(k+1):(t-1)} | S_k)$ right-to-left.

Base case ($k = t - 1$, no future observations)

$$b_{t:(t-1)} = \langle 1, 1 \rangle$$


Recursive case ($k < t - 1$)

$$b_{(k+1):(t-1)} = \sum_{s_{k+1}} P(o_{k+1} | s_{k+1}) b_{(k+2):(t-1)} P(s_{k+1} | S_k)$$

Same factor sizes as forward recursion — total cost $O(t)$ for the whole backward pass.

No normalization here: b is a likelihood, not a probability. Normalize *after* multiplying $f \times b$.

A smoothing example: setup



Umbrella story, given $O_0 = O_1 = O_2 = t$. Compute the *smoothed* beliefs:

- $P(S_0 | o_{0:2})$: did it rain on **day 0**?
- $P(S_1 | o_{0:2})$: did it rain on **day 1**?

From L10 we already have $f_{0:0} = \langle 0.818, 0.182 \rangle$ and $f_{0:1} = \langle 0.883, 0.117 \rangle$.

Backward pass: compute $b_{2:2}$, then $b_{1:2}$

$$b_{2:2} = P(o_2|S_1) = \sum_{s_2} P(o_2|s_2) b_{3:2} P(s_2|S_1) \text{ with } b_{3:2} = \langle 1, 1 \rangle:$$

$$= 0.9 \cdot 1 \cdot \langle 0.7, 0.3 \rangle + 0.2 \cdot 1 \cdot \langle 0.3, 0.7 \rangle$$

$$= \langle 0.63, 0.27 \rangle + \langle 0.06, 0.14 \rangle = \langle \mathbf{0.69}, \mathbf{0.41} \rangle$$

$$b_{1:2} = P(o_{1:2}|S_0) = \sum_{s_1} P(o_1|s_1) b_{2:2} P(s_1|S_0):$$

$$= 0.9 \cdot 0.69 \cdot \langle 0.7, 0.3 \rangle + 0.2 \cdot 0.41 \cdot \langle 0.3, 0.7 \rangle$$

$$= \langle 0.4347, 0.1863 \rangle + \langle 0.0246, 0.0574 \rangle = \langle \mathbf{0.4593}, \mathbf{0.2437} \rangle$$

Combine: $P(S_k | o_{0:2}) = \alpha f_{0:k} b_{(k+1):2}$

$$P(S_1 | o_{0:2}) = \alpha f_{0:1} \cdot b_{2:2}:$$

$$= \alpha \langle 0.883, 0.117 \rangle \cdot \langle 0.69, 0.41 \rangle = \alpha \langle 0.6093, 0.0480 \rangle = \langle \mathbf{0.927}, \mathbf{0.073} \rangle$$

$$P(S_0 | o_{0:2}) = \alpha f_{0:0} \cdot b_{1:2}:$$

$$= \alpha \langle 0.818, 0.182 \rangle \cdot \langle 0.4593, 0.2437 \rangle = \alpha \langle 0.3757, 0.0444 \rangle = \langle \mathbf{0.894}, \mathbf{0.106} \rangle$$

$$P(S_1 = T | \dots)$$

Filtering 0.883
with $o_{0:1}$
only

Smoothing **0.927**
with $o_{0:2}$

The extra "umbrella on day 2" makes "rain on day 1" more likely — future evidence helps.

Deriving the backward recursion


Five steps from $P(o_{(k+1):(t-1)} | S_k)$ to the recursive form:

$$\begin{aligned} & P(o_{(k+1):(t-1)} | S_k) \\ &= \sum_{s_{k+1}} P(o_{(k+1):(t-1)} \wedge s_{k+1} | S_k) && \text{sum rule} \\ &= \sum_{s_{k+1}} P(o_{(k+1):(t-1)} | s_{k+1} \wedge S_k) P(s_{k+1} | S_k) && \text{chain rule} \\ &= \sum_{s_{k+1}} P(o_{(k+1):(t-1)} | s_{k+1}) P(s_{k+1} | S_k) && \text{Markov} \\ &= \sum_{s_{k+1}} P(o_{k+1} \wedge o_{(k+2):(t-1)} | s_{k+1}) P(s_{k+1} | S_k) && \text{rewrite} \\ &= \sum_{s_{k+1}} P(o_{k+1} | s_{k+1}) P(o_{(k+2):(t-1)} | s_{k+1}) P(s_{k+1} | S_k) && \text{Markov} \end{aligned}$$

The final factor is exactly $P(o_{k+1} | s_{k+1}) b_{(k+2):(t-1)} P(s_{k+1} | S_k)$ — the backward recursion.

The forward-backward algorithm

Compute the smoothed posterior at every time step k with one forward sweep + one backward sweep.




Combine the two sweeps at every k

$$P(S_k \mid o_{0:(t-1)}) = \alpha f_{0:k} \cdot b_{(k+1):(t-1)}$$

Forward gives every $f_{0:k}$ (recap of L10), backward gives every $b_{(k+1):(t-1)}$ — total cost $O(t)$, one sweep each direction.

Try it: smoothing on a 4-step HMM

Compute $P(S_2 \mid o_0 \wedge o_1 \wedge o_2 \wedge \neg o_3)$ — here $k = 2$, $t = 4$.



Numbers

$$P(s_0) = 0.4$$

$$P(s_t | s_{t-1}) = 0.7$$

$$P(s_t | \neg s_{t-1}) = 0.2$$

$$P(o_t | s_t) = 0.9$$

$$P(o_t | \neg s_t) = 0.2$$

Forward (from L10): $f_{0:2} \approx \langle 0.884, 0.116 \rangle$

Backward: $b_{3:3} = 0.1 \langle 0.7, 0.2 \rangle + 0.8 \langle 0.3, 0.8 \rangle = \langle 0.31, 0.66 \rangle$

Smoothed: $\alpha \langle 0.884, 0.116 \rangle \cdot \langle 0.31, 0.66 \rangle = \alpha \langle 0.274, 0.0766 \rangle = \boxed{\langle 0.782, 0.218 \rangle}$

The most-likely explanation

Find the single **best sequence of hidden states** that explains the observations:

$$\hat{s}_{0:t-1} = \arg \max_{s_{0:t-1}} P(s_{0:t-1} | o_{0:t-1})$$

- Generalize: $S_t \in \{0, 1, \dots, n - 1\}$ — not just binary.
- Transition matrix $A \in \mathbb{R}^{n \times n}$, emission matrix $O \in \mathbb{R}^{n \times m}$.

Brute force: enumerate every state sequence.


n^t sequences — intractable.

Dynamic programming to the rescue: the **Viterbi algorithm**.

Dynamic programming: define $\pi_k(j)$

$$\pi_k(j) = \max_{s_{0:(k-1)}} P(s_{0:(k-1)}, S_k = j \mid o_{0:k})$$

The highest probability of any state sequence $s_{0:k}$ ending with $S_k = j$, given observations $o_{0:k}$.



Trellis: one column per time step, one node per state value. Red = best path entering each node.

Viterbi recursion

Base case

$$\pi_0(j) = \alpha P(S_0 = j) P(o_0 | S_0 = j)$$

Recursive case

$$\pi_k(j) = P(o_k | S_k = j) \cdot \max_z \left[\pi_{k-1}(z) P(S_k = j | S_{k-1} = z) \right]$$

Backpointer

$$\phi_k(j) = \arg \max_z \left[\pi_{k-1}(z) P(S_k = j | S_{k-1} = z) \right]$$

$\pi_k(j)$ = best probability of any path ending at state j at time k . $\phi_k(j)$ remembers which predecessor produced it — we'll use it to trace the best sequence back.


Viterbi: the algorithm

viterbi(observations, transition, emission)

1. Initialize $\pi_0(j) = \alpha P(S_0 = j) P(o_0|S_0 = j)$ for each state j ; $\phi_0(j) = \text{none}$.
2. **Forward sweep:** for $k = 1, \dots, t - 1$ and each j , compute $\pi_k(j)$ and record the argmax in $\phi_k(j)$.
3. **Best terminal:** $\hat{s}_{t-1} = \arg \max_j \pi_{t-1}(j)$.
4. **Traceback:** for $k = t - 1, \dots, 1$, set $\hat{s}_{k-1} = \phi_k(\hat{s}_k)$.
5. Return the sequence $(\hat{s}_0, \hat{s}_1, \dots, \hat{s}_{t-1})$.

One sweep fills the trellis; one walk back along the backpointers extracts the MAP sequence.

Viterbi: trellis walk-through + complexity



Best path: $\hat{s} = (1, 0, 1, 2, 1)$.

Complexity: $O(t \cdot n^2)$ — one trellis edge per cell per step. Linear in time, quadratic in #states.

Learning goals (recap)

- ✓ Compute the **smoothing** probability for a time step in an HMM.
- ✓ **Justify** each step in the smoothing & backward derivations.
- ✓ Describe the **forward-backward** algorithm.
- ✓ Describe the **Viterbi** algorithm.

Search ›

Uncertainty ›

Decisions ›

Learning

Next module: Decision Making

So far we have *beliefs*. Next we add *preferences* — how should an agent **act** when outcomes are uncertain?

- **L12** — decision theory: utility, MEU, decision networks, value of information.
- **L13–L14** — Markov decision processes: Bellman equation, value & policy iteration.
- **L15** — reinforcement learning: ADP, Q-learning, SARSA.