

# CS 486/686

## Hidden Markov Models II

Yuntian Deng

---

Lecture 11

RN 14.2.2

# Reminders



**Chat 4 is due TODAY** (Tue Jun 16); Chat 5 due Thu Jun 18 — both 11:59 PM on Chrysalis.



**Assignment 1 deadline extended to Thu Jun 25** (was Jun 18). The PDF, starter code, and model file were **updated on LEARN** — please re-download.



**Assignment 1 questions → Dake Zhang on Piazza.**

Search ›

Uncertainty ›

Decisions ›

Learning


## Learning goals

- Compute the **smoothing** probability for a time step in an HMM.
- **Justify** each step in the derivation of the smoothing formula.
- Describe the **forward-backward** algorithm.
- Describe the **Viterbi** algorithm.

# Recap: filtering (from L10)

From past evidence  $o_{0:k}$ , filtering gives the posterior over the *current* state:

$$f_{0:k} = P(S_k | o_{0:k}) = \alpha P(o_k | S_k) \sum_{s_{k-1}} P(S_k | s_{k-1}) f_{0:(k-1)}$$




- Yellow = query state  $S_k$  (the *present*).
- Gray = observed evidence.
- Cost:  $O(k)$  ops via the recursive predict-update.

Today: a query about a **past** state, given all the evidence we have.

# Smoothing: query a past state

$$P(S_k | o_{0:(t-1)}), \quad 0 \leq k \leq t - 1$$



## Filtering (L10)

$P(S_k | o_{0:k})$  — the present state from past evidence.

## Smoothing (today)

$P(S_k | o_{0:(t-1)}), k < t - 1$  — a past state from all evidence (past *and* future).

# Deriving the smoothing formula

Four small steps from  $P(S_k | o_{0:(t-1)})$  to the forward-backward product:

$$P(S_k | o_{0:(t-1)})$$

$$= P(S_k | o_{(k+1):(t-1)} \wedge o_{0:k})$$

*rewrite*

$$= \alpha P(S_k | o_{0:k}) P(o_{(k+1):(t-1)} | S_k \wedge o_{0:k})$$

*Bayes' rule*

$$= \alpha P(S_k | o_{0:k}) P(o_{(k+1):(t-1)} | S_k)$$

*Markov*

$$= \alpha f_{0:k} b_{(k+1):(t-1)}$$


*definitions*

Try to identify the rule before the label appears.

# Smoothing = forward × backward

Split the evidence around  $S_k$  into past and future:

$$P(S_k | o_{0:(t-1)}) = \alpha \underbrace{P(S_k | o_{0:k})}_{\text{forward } f_{0:k}} \underbrace{P(o_{(k+1):(t-1)} | S_k)_{\text{backward } b_{(k+1):(t-1)}}$$



Two sweeps meet at  $S_k$ ; multiply pointwise and normalize.

# Backward recursion

Compute  $b_{(k+1):(t-1)} \triangleq P(o_{(k+1):(t-1)} | S_k)$  right-to-left.

**Base case ( $k = t - 1$ , no future observations)**

$$b_{t:(t-1)} = \langle 1, 1 \rangle$$

**Recursive case ( $k < t - 1$ )**


$$b_{(k+1):(t-1)} = \sum_{s_{k+1}} P(o_{k+1} | s_{k+1}) b_{(k+2):(t-1)} P(s_{k+1} | S_k)$$

Same factor sizes as forward recursion — total cost  $O(t)$  for the whole backward pass.

No normalization here:  $b$  is a likelihood, not a probability. Normalize *after* multiplying  $f \times b$ .



# A smoothing example: setup



Umbrella story, given  $O_0 = O_1 = O_2 = t$ . Compute the *smoothed* beliefs:

- $P(S_0 | o_{0:2})$ : did it rain on **day 0**?
- $P(S_1 | o_{0:2})$ : did it rain on **day 1**?

From L10 we already have  $f_{0:0} = \langle 0.818, 0.182 \rangle$  and  $f_{0:1} = \langle 0.883, 0.117 \rangle$ .

## Backward pass: compute $b_{2:2}$ , then $b_{1:2}$

$$b_{2:2} = P(o_2|S_1) = \sum_{s_2} P(o_2|s_2) b_{3:2} P(s_2|S_1) \text{ with } b_{3:2} = \langle 1, 1 \rangle:$$

$$= 0.9 \cdot 1 \cdot \langle 0.7, 0.3 \rangle + 0.2 \cdot 1 \cdot \langle 0.3, 0.7 \rangle$$

$$= \langle 0.63, 0.27 \rangle + \langle 0.06, 0.14 \rangle = \langle \mathbf{0.69}, \mathbf{0.41} \rangle$$

$$b_{1:2} = P(o_{1:2}|S_0) = \sum_{s_1} P(o_1|s_1) b_{2:2} P(s_1|S_0):$$

$$= 0.9 \cdot 0.69 \cdot \langle 0.7, 0.3 \rangle + 0.2 \cdot 0.41 \cdot \langle 0.3, 0.7 \rangle$$

$$= \langle 0.4347, 0.1863 \rangle + \langle 0.0246, 0.0574 \rangle = \langle \mathbf{0.4593}, \mathbf{0.2437} \rangle$$

**Combine:**  $P(S_k | o_{0:2}) = \alpha f_{0:k} b_{(k+1):2}$

$$P(S_1 | o_{0:2}) = \alpha f_{0:1} \cdot b_{2:2}:$$

$$= \alpha \langle 0.883, 0.117 \rangle \cdot \langle 0.69, 0.41 \rangle = \alpha \langle 0.6093, 0.0480 \rangle = \langle \mathbf{0.927}, \mathbf{0.073} \rangle$$

$$P(S_0 | o_{0:2}) = \alpha f_{0:0} \cdot b_{1:2}:$$

$$= \alpha \langle 0.818, 0.182 \rangle \cdot \langle 0.4593, 0.2437 \rangle = \alpha \langle 0.3757, 0.0444 \rangle = \langle \mathbf{0.894}, \mathbf{0.106} \rangle$$

$$P(S_1 = T | \dots)$$

Filtering      0.883  
with  $o_{0:1}$   
only

---

Smoothing    **0.927**  
with  $o_{0:2}$

The extra "umbrella on day 2" makes "rain on day 1" more likely — future evidence helps.

# Deriving the backward recursion


Five steps from  $P(o_{(k+1):(t-1)} | S_k)$  to the recursive form:

$$\begin{aligned} &P(o_{(k+1):(t-1)} | S_k) \\ &= \sum_{s_{k+1}} P(o_{(k+1):(t-1)} \wedge s_{k+1} | S_k) && \text{sum rule} \\ &= \sum_{s_{k+1}} P(o_{(k+1):(t-1)} | s_{k+1} \wedge S_k) P(s_{k+1} | S_k) && \text{chain rule} \\ &= \sum_{s_{k+1}} P(o_{(k+1):(t-1)} | s_{k+1}) P(s_{k+1} | S_k) && \text{Markov} \\ &= \sum_{s_{k+1}} P(o_{k+1} \wedge o_{(k+2):(t-1)} | s_{k+1}) P(s_{k+1} | S_k) && \text{rewrite} \\ &= \sum_{s_{k+1}} P(o_{k+1} | s_{k+1}) P(o_{(k+2):(t-1)} | s_{k+1}) P(s_{k+1} | S_k) && \text{Markov} \end{aligned}$$

The final factor is exactly  $P(o_{k+1} | s_{k+1}) b_{(k+2):(t-1)} P(s_{k+1} | S_k)$  — the backward recursion.

# The forward-backward algorithm

Compute the smoothed posterior at every time step  $k$  with one forward sweep + one backward sweep.




**Combine the two sweeps at every  $k$**

$$P(S_k \mid o_{0:(t-1)}) = \alpha f_{0:k} \cdot b_{(k+1):(t-1)}$$

Forward gives every  $f_{0:k}$  (recap of L10), backward gives every  $b_{(k+1):(t-1)}$  — total cost  $O(t)$ , one sweep each direction.

# Try it: smoothing on a 4-step HMM

Compute  $P(S_2 \mid o_0 \wedge o_1 \wedge o_2 \wedge \neg o_3)$  — here  $k = 2$ ,  $t = 4$ .



## Numbers

$$P(s_0) = 0.4$$

$$P(s_t | s_{t-1}) = 0.7$$

$$P(s_t | \neg s_{t-1}) = 0.2$$

$$P(o_t | s_t) = 0.9$$

$$P(o_t | \neg s_t) = 0.2$$

Forward (from L10):  $f_{0:2} \approx \langle 0.884, 0.116 \rangle$

Backward:  $b_{3:3} = 0.1 \langle 0.7, 0.2 \rangle + 0.8 \langle 0.3, 0.8 \rangle = \langle 0.31, 0.66 \rangle$

Smoothed:  $\alpha \langle 0.884, 0.116 \rangle \cdot \langle 0.31, 0.66 \rangle = \alpha \langle 0.274, 0.0766 \rangle = \boxed{\langle 0.782, 0.218 \rangle}$

# The most-likely explanation

Find the single **best sequence of hidden states** that explains the observations:

$$\hat{s}_{0:t-1} = \arg \max_{s_{0:t-1}} P(s_{0:t-1} | o_{0:t-1})$$

- Generalize:  $S_t \in \{0, 1, \dots, n - 1\}$  — not just binary.
- Transition matrix  $A \in \mathbb{R}^{n \times n}$ , emission matrix  $O \in \mathbb{R}^{n \times m}$ .

**Brute force:** enumerate every state sequence.


$n^t$  sequences — intractable.

Dynamic programming to the rescue: the **Viterbi algorithm**.

# Dynamic programming: define $\pi_k(j)$

$$\pi_k(j) = \max_{s_{0:(k-1)}} P(s_{0:(k-1)}, S_k = j \mid o_{0:k})$$

The highest probability of any state sequence  $s_{0:k}$  ending with  $S_k = j$ , given observations  $o_{0:k}$ .



Trellis: one column per time step, one node per state value. Red = best path entering each node.



# Why the recursion holds

The best path to state  $j$  at time  $k$  reuses the best path to some predecessor  $z$  at  $k - 1$

:

$$\pi_k(j) = \max_{s_{0:(k-1)}} P(s_{0:(k-1)}, S_k=j \mid o_{0:k})$$

$$\propto \max_{s_{0:(k-1)}} P(s_{0:(k-1)}, S_k=j, o_{0:k}) \quad \text{drop } P(o_{0:k})$$

$$= \max_{s_{0:(k-1)}} P(s_{0:(k-1)}, o_{0:(k-1)}) \underbrace{P(S_k=j \mid s_{k-1})}_{\text{transition}} \underbrace{P(o_k \mid j)}_{\text{emission}} \quad \text{HMM factorization}$$

$$= P(o_k \mid j) \max_{s_{k-1}} \max_{s_{0:(k-2)}} P(s_{0:(k-2)}, s_{k-1}, o_{0:(k-1)}) P(j \mid s_{k-1}) \quad \text{nest the prefix max}$$

$$= P(o_k \mid j) \max_z P(j \mid z) \max_{s_{0:(k-2)}} P(s_{0:(k-2)}, S_{k-1}=z, o_{0:(k-1)}) \quad \text{pull out } P(j \mid z)$$

$$= P(o_k \mid j) \max_z [P(j \mid z) \pi_{k-1}(z)] \quad \text{inner max} = \pi_{k-1}(z)$$

The inner max has no  $j$  in it — so **one best path per node** serves every successor ( $\phi_k(j)$  records the winning  $z$ ). Constants don't change the argmax, so we work with the joint.

# Viterbi recursion

## Base case

$$\pi_0(j) = \alpha P(S_0 = j) P(o_0 | S_0 = j)$$

## Recursive case

$$\pi_k(j) = P(o_k | S_k = j) \cdot \max_z \left[ \pi_{k-1}(z) P(S_k = j | S_{k-1} = z) \right]$$

## Backpointer

$$\phi_k(j) = \arg \max_z \left[ \pi_{k-1}(z) P(S_k = j | S_{k-1} = z) \right]$$

$\pi_k(j)$  = best probability of any path ending at state  $j$  at time  $k$ .  $\phi_k(j)$  remembers which predecessor produced it — we'll use it to trace the best sequence back.

# Viterbi: the algorithm

## **viterbi(observations, transition, emission)**

1. Initialize  $\pi_0(j) = \alpha P(S_0 = j) P(o_0|S_0 = j)$  for each state  $j$ ;  $\phi_0(j) = \text{none}$ .
2. **Forward sweep:** for  $k = 1, \dots, t - 1$  and each  $j$ , compute  $\pi_k(j)$  and record the argmax in  $\phi_k(j)$ .
3. **Best terminal:**  $\hat{s}_{t-1} = \arg \max_j \pi_{t-1}(j)$ .
4. **Traceback:** for  $k = t - 1, \dots, 1$ , set  $\hat{s}_{k-1} = \phi_k(\hat{s}_k)$ .
5. Return the sequence  $(\hat{s}_0, \hat{s}_1, \dots, \hat{s}_{t-1})$ .

One sweep fills the trellis; one walk back along the backpointers extracts the MAP sequence.

# Why Viterbi returns the true argmax

Recall  $\pi_k(j) = \max_{s_{0:(k-1)}} P(s_{0:(k-1)}, S_k=j, o_{0:k})$  and  $\phi_k(j) =$

$\arg \max_z [\pi_{k-1}(z) P(j \mid z)]$ . Goal:  $\hat{s}_{0:t-1} = \arg \max_{s_{0:t-1}} P(s_{0:t-1}, o_{0:t-1})$ .

$$\hat{s}_{t-1}, \hat{s}_{t-2}, \dots, \hat{s}_0 = \arg \max_{s_{0:t-1}} P(s_{0:t-1}, o_{0:t-1})$$

*the MAP path*

$$= \arg \max_{s_{t-1}} \arg \max_{s_{0:t-2}} P(s_{0:t-1}, o_{0:t-1})$$

*split: pick  $s_{t-1}$ , then the rest*

$$\hat{s}_{t-1} = \arg \max_{s_{t-1}} \max_{s_{0:t-2}} P(s_{0:t-1}, o_{0:t-1}) = \arg \max_{s_{t-1}} \pi_{t-1}(s_{t-1})$$

*outer argmax = best final state*

$$\hat{s}_{t-2} = \arg \max_{s_{t-2}} \pi_{t-2}(s_{t-2}) P(\hat{s}_{t-1} \mid s_{t-2}) = \phi_{t-1}(\hat{s}_{t-1})$$


*extract one more (drop const. emission)*

$$\hat{s}_{k-1} = \phi_k(\hat{s}_k), \quad k = t-1, \dots, 1$$

*repeat down the chain* ■

Splitting the argmax one state at a time turns the global search into  $t$  tiny argmaxes — and each one is exactly a backpointer  $\phi_k$ .

# Viterbi: trellis walk-through + complexity



Best path:  $\hat{s} = (1, 0, 1, 2, 1)$ .

Complexity:  $O(t \cdot n^2)$  — one trellis edge per cell per step. Linear in time, quadratic in #states.

## Learning goals (recap)

- ✓ Compute the **smoothing** probability for a time step in an HMM.
- ✓ **Justify** each step in the smoothing & backward derivations.
- ✓ Describe the **forward-backward** algorithm.
- ✓ Describe the **Viterbi** algorithm.

Search ›

Uncertainty ›

**Decisions** ›

Learning

## Next module: Decision Making

So far we have *beliefs*. Next we add *preferences* — how should an agent **act** when outcomes are uncertain?

- **L12** — decision theory: utility, MEU, decision networks, value of information.
- **L13–L14** — Markov decision processes: Bellman equation, value & policy iteration.
- **L15** — reinforcement learning: ADP, Q-learning, SARSA.