

# CS 486/686

## Value & Policy Iteration

Yuntian Deng

---

Lecture 14

RN 17.2 · PM 9.5.2–9.5.3



## Learning goals

- Trace + implement the **value iteration** algorithm for solving an MDP.
- Trace + implement the **policy iteration** algorithm for solving an MDP.

# Recap: $\pi^\star = \arg \max_a Q^\star$

From L13, the optimal policy is one greedy step from  $V^\star$ :

$$\pi^\star(s) = \arg \max_a Q^\star(s, a), \quad Q^\star(s, a) = R(s) + \gamma \sum_{s'} P(s' | s, a) V^\star(s').$$

## Today's question

How do we actually *compute*  $V^\star$ ?

Two algorithms: **value iteration** (iterate values) and **policy iteration** (iterate policies).

# Value functions

- $V^\pi(s)$  — value of being in state  $s$  under policy  $\pi$ .
- $V^*(s)$  — value of  $s$  under the **optimal** policy  $\pi^*$ .
- $Q^\pi(s, a)$  — value of taking action  $a$  in  $s$ , then following  $\pi$ .
- $Q^*(s, a)$  — same, but under  $\pi^*$ .
- $\pi(a \mid s)$  — the policy, a distribution over actions per state.

"Value" = expected discounted future reward.

# Formal definitions

**Return** (discounted future reward from time  $t$ ):

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = R_{t+1} + \gamma G_{t+1}.$$

**Value functions are expected returns:**

$$V^\pi(s) = \mathbb{E}_\pi [G_t \mid s_t = s], \quad Q^\pi(s, a) = \mathbb{E}_\pi [G_t \mid s_t = s, a_t = a].$$


**The two are related** (each is a sum of the other):

$$V^\pi(s) = \sum_a \pi(a \mid s) Q^\pi(s, a)$$

$$Q^\pi(s, a) = R(s) + \gamma \sum_{s'} P(s' \mid s, a) V^\pi(s')$$

# The backup diagram

A picture of the V/Q recursion: **policy** branches on actions; **environment** branches on next states.



# The Bellman optimality equation

For the optimal policy, choose the *best* action at every state:

$$V^*(s) = \max_a Q^*(s, a), \quad Q^*(s, a) = R(s) + \gamma \sum_{s'} P(s' \mid s, a) V^*(s').$$

Substitute one into the other:

**Bellman equation**

$$V^*(s) = R(s) + \gamma \max_a \sum_{s'} P(s' \mid s, a) V^*(s').$$

$V^*$  is the **unique** solution to this system of  $|S|$  equations.

# Apply Bellman to $V^*(s_{11})$

|          |     |     |      |
|----------|-----|-----|------|
| $s_{11}$ | 0.0 | 0.0 | 0.04 |
| 0.0      | X   | 0.0 | -1   |
| 0.0      | 0.0 | 0.0 | +1   |

From  $s_{11}$ , each action gives a different mixture (0.8 intended + 0.1 each side):

$$V^*(s_{11}) = -0.04 + \gamma \max \left\{ \begin{array}{l} 0.8 V^*(s_{12}) + 0.1 V^*(s_{21}) + 0.1 V^*(s_{11}), \text{ // right} \\ 0.9 V^*(s_{11}) + 0.1 V^*(s_{12}), \text{ // up} \\ 0.9 V^*(s_{11}) + 0.1 V^*(s_{21}), \text{ // left} \\ 0.8 V^*(s_{21}) + 0.1 V^*(s_{12}) + 0.1 V^*(s_{11}) \text{ // down} \end{array} \right\}$$

**Q1.** Can we solve this system of Bellman equations in polynomial time?

**No.** The system is *nonlinear* because of the **max** — there is no general efficient solver. We'll use an **iterative** algorithm instead.



# The value iteration algorithm

Treat the Bellman equation as an *update rule*. Let  $V_i(s)$  be the  $i^{\text{th}}$  estimate.

value-iteration(MDP,  $\epsilon$ )

1. Initialize  $V_0(s)$  arbitrarily (e.g. 0).
2. For  $i = 0, 1, 2, \dots$ : for each state  $s$ , apply the Bellman update

$$V_{i+1}(s) \leftarrow R(s) + \gamma \max_a \sum_{s'} P(s' \mid s, a) V_i(s').$$

3. Stop when  $\max_s |V_{i+1}(s) - V_i(s)| < \epsilon$ .

**Guarantee:**  $V_i \rightarrow V^*$  as  $i \rightarrow \infty$ .

- **Synchronous:** keep  $V_i$  frozen while computing all of  $V_{i+1}$ .
- **Asynchronous:** update entries in place, in any order. Often converges faster

# Apply VI to the grid

**Setup:**  $\gamma = 1$ ,  $R(s) = -0.04$  for non-goal states. Terminals:  $V(s_{24}) = -1$ ,  $V(s_{34}) = +1$ . Init  $V_0(s) = 0$  for all non-terminals.

$V_0(s)$

|   |   |   |    |
|---|---|---|----|
| 0 | 0 | 0 | 0  |
| 0 | X | 0 | -1 |
| 0 | 0 | 0 | +1 |

After one Bellman update, only the cells adjacent to a goal will see something different from  $-0.04$ .

# One iteration: $V_1(s)$

Apply  $V_1(s) \leftarrow R(s) + \gamma \max_a \sum_{s'} P(s' | s, a) V_0(s')$  at every cell. Two cells worth a closer look:

$V_1(s)$

|       |      |       |      |
|-------|------|-------|------|
| -0.00 | 0.00 | 0.00  | 0.04 |
| -0.00 | X    | -0.00 | -1   |
| -0.00 | 0.00 | 0.76  | +1   |

Q2.  $V_1(s_{23})$  (yellow): all neighbors are 0; best action is Left (avoids  $-1$ ).  $\Rightarrow -0.04 + 0 = \boxed{-0.04}$ .

Q3.  $V_1(s_{33})$  (green): Right has  $0.8 \cdot 1 + 0.1 \cdot 0 + 0.1 \cdot 0 = 0.8$ .  $\Rightarrow -0.04 + 0.8 = \boxed{0.76}$ .

The +1 reward has "leaked" one cell up; the  $-1$  hasn't propagated because we avoid it.

# Two iterations: $V_2(s)$

Apply Bellman again, using  $V_1$  as input. Three cells worth a closer look:

$V_2(s)$

|       |      |      |      |
|-------|------|------|------|
| -0.00 | 0.00 | 0.00 | 0.08 |
| -0.00 | X    | 0.46 | -1   |
| -0.00 | 0.50 | 0.83 | +1   |

Q4.  $V_2(s_{33})$  : Right gives  $0.8(+1) + 0.1(-0.04) + 0.1(0.76) = 0.872. \Rightarrow$   
 $-0.04 + 0.872 = \boxed{0.832}.$

Q5.  $V_2(s_{23})$  : Down gives  $0.8(0.76) + 0.1(-0.04) + 0.1(-1) = 0.504. \Rightarrow$   
 $-0.04 + 0.504 = \boxed{0.464}.$

Q6.  $V_2(s_{32})$  : Right gives  $0.8(0.76) + 0.1(-0.04) + 0.1(-0.04) = 0.6. \Rightarrow$   
 $-0.04 + 0.6 = \boxed{0.56}.$

Values propagate one cell per iteration. After enough sweeps they converge to the L13  $V^*$ .

# Policy iteration

**Insight:** if one action is clearly better, you don't need precise utility values — just enough to rank actions.

policy-iteration(MDP)

1. Start from an arbitrary policy  $\pi_0$ .
2. Repeat:
  - **Policy evaluation:** compute  $V^{\pi_i}(s)$ , the utility of every state under  $\pi_i$ .
  - **Policy improvement:**  $\pi_{i+1}(s) = \arg \max_a \sum_{s'} P(s' \mid s, a) V^{\pi_i}(s')$ .
3. Until  $\pi_{i+1} = \pi_i$ .

Terminates in finitely many steps: there are finitely many policies and each iteration strictly improves.

# Policy evaluation vs Bellman: the max disappears

In PE the action is *fixed* by  $\pi$ , so the system becomes linear in  $V$ :

## Policy evaluation (linear)

$$V^\pi(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) V^\pi(s')$$

No **max** — the action is fixed.  $|S|$   
linear equations in  $|S|$  unknowns.

linear: easy to solve

## Bellman optimality (nonlinear)

$$V^*(s) = R(s) + \gamma \max_a \sum_{s'} P(s' | s, a) V^*(s')$$

The **max** makes it nonlinear — need  
an iterative algorithm (VI).

nonlinear: harder

Policy improvement reintroduces the **arg max**:  $\pi_{i+1}(s) = \arg \max_a \sum_{s'} P(s' | s, a) V^{\pi_i}(s')$ .

# How do we perform policy evaluation?

Two ways to solve the linear system  $V^\pi(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) V^\pi(s')$ :

## Exactly

Linear algebra: a single solve of an  $|S| \times |S|$  system.

**cost:**  $O(|S|^3)$

## Iteratively (cheap)

Run  $m$  simplified Bellman updates (no **max**):

$$V_{j+1}(s) \leftarrow R(s) + \gamma \sum_{s'} P(s' | s, \pi(s)) V_j(s')$$

**cost:**  $O(m \cdot |S| \cdot |A_{\text{avg}}|)$

In practice, a small  $m$  is usually enough — PI converges in far fewer outer iterations than

## Learning goals (recap)

- ✓ Trace and implement **value iteration** for solving an MDP.
- ✓ Trace and implement **policy iteration** for solving an MDP.



## Next: reinforcement learning

Today we knew the dynamics  $P$  and rewards  $R$  of the MDP. What if we *don't*?

L15: learn the policy by **interacting with the environment**.