

# CS 486/686

# Reinforcement Learning

Yuntian Deng

---

Lecture 15

RN 22.1–22.3 · PM 12.1, 12.5, 12.8

# Heads-up: next week is asynchronous

 I'll be away at ICML the week of July 7. **L17 (Tue Jul 7)** and **L18 (Thu Jul 9)** will be **pre-recorded videos** — no in-person class those days. Watch on your own time (links posted before class).

 **Deadlines are unchanged:** Chat 8 + the CS686 project proposal are due **Tue Jul 7**; Chat 9 is out and **Assignment 2 is due Thu Jul 9**.

 Questions? Post on **Piazza** — the TAs are available all week.

Search ›

Uncertainty ›


**Decisions** ›

Learning

## Learning goals

- Trace + implement **passive adaptive dynamic programming** (passive ADP).
- Explain the **exploration vs exploitation** trade-off.
- Trace + implement **active ADP**.
- Trace + implement **Q-learning**.
- Trace + implement **SARSA**.
- Tell **on-policy** from **off-policy** learning.

# The RL setting: no $P$ , no $R$ up front



An MDP whose **transition  $P$**  and **reward  $R$**  are unknown. The agent sees state  $s$  and reward  $r$ , picks an action  $a$ , and observes the next state  $s'$ .

**Goal:** maximize discounted return — while learning the world as it goes.

# Passive learning: just evaluate a fixed policy

Fix a policy  $\pi$ ; goal: learn  $V^\pi(s)$  for every state  $s$ . **The catch:** we no longer know the reward  $R(s)$  or the transitions  $P(s' \mid s, a)$  — we only get to *observe* experiences.

## Adaptive Dynamic Programming (ADP)

1. **Learn** the transition  $P(s' \mid s, a)$  and reward  $R(s)$  from observed experience.
2. **Solve** the Bellman policy-evaluation equations:

$$V^\pi(s) = R(s) + \gamma \sum_{s'} P(s' \mid s, \pi(s)) V^\pi(s').$$

A **model-based** approach — it builds an explicit model of the world.

# The passive ADP algorithm

## passive-ADP( $\pi$ )

1. Initialize  $R(s)$ ,  $N(s, a)$ ,  $N(s, a, s')$ ,  $V^\pi(s)$ .
2. Follow  $\pi$  and observe an experience  $\langle s, a, s', r \rangle$ .
3. Update reward:  $R(s) \leftarrow r$ .
4. Update counts and transition probability:

$$N(s, a) += 1, \quad N(s, a, s') += 1, \quad P(s' \mid s, a) = \frac{N(s, a, s')}{N(s, a)}.$$

5. Re-solve Bellman to get  $V^\pi(s)$  (exactly or iteratively).

# Passive ADP: update the model

Same grid world as before (the agent moves in the intended direction, or slips to either side) — but now **we don't know** those slip probabilities, so we **count outcomes** to estimate  $P$ .

|          |    |
|----------|----|
| $s_{11}$ | +1 |
| $s_{21}$ | -1 |

$\pi(s_{11}) = \text{down}, \pi(s_{21}) = \text{right}, \gamma = 0.9.$

Counts so far:  $N(s, a) = 5$ , with  $N(s, a, s') = 3$  for the intended direction and 1 for each side.

New experience:  $\langle s_{11}, \text{down}, s_{21}, -0.04 \rangle$ .

Increment counts:  $N(s_{11}, \text{down}): 5 \rightarrow 6$ ,  
 $N(s_{11}, \text{down}, s_{21}): 3 \rightarrow 4$ .

Re-estimate  $P(\cdot \mid s_{11}, \text{down})$  as count fractions:

$$s_{21}: \frac{4}{6}, \quad \text{stay } s_{11}: \frac{1}{6}, \quad +1: \frac{1}{6}$$

(Going down can slip sideways: left bumps the wall  $\rightarrow$  stay; right  $\rightarrow$  the +1 cell.)

# Passive ADP: re-solve with the new model

Reminder — policy evaluation:  $V^\pi(s) = R(s) + \gamma \sum_{s'} P(s' \mid s, \pi(s)) V^\pi(s')$  ( $R = -0.04, \gamma = 0.9$ ).

Estimate  $P(\cdot \mid s_{21}, \text{right})$  the same way (counts 3 intended, 1 each side, out of 5): right  $\rightarrow -1$  is  $\frac{3}{5}$ ; up  $\rightarrow s_{11}$  is  $\frac{1}{5}$ ; down bumps the wall  $\rightarrow$  stay  $s_{21}$  is  $\frac{1}{5}$ .

$$V(s_{11}) = -0.04 + 0.9 \left[ \frac{4}{6} V(s_{21}) + \frac{1}{6} V(s_{11}) + \frac{1}{6} (+1) \right]$$

$$V(s_{21}) = -0.04 + 0.9 \left[ \frac{3}{5} (-1) + \frac{1}{5} V(s_{11}) + \frac{1}{5} V(s_{21}) \right]$$

Two linear equations, two unknowns — solve:

$$V(s_{11}) = -0.438, \quad V(s_{21}) = -0.803.$$



## Quick check: probability update

**Q1.** In a 3x4 grid, suppose  $N(s_{23}, \text{down}) = 23$  and  $N(s_{23}, \text{down}, s_{24}) = 2$ . The agent tries to move *down* but slips and lands in  $s_{24}$ . What is the updated  $P(s_{24} \mid s_{23}, \text{down})$ ?

A.  $2/23 \approx 0.087$

B.  $3/23 \approx 0.130$

C.  $3/24 = 0.125$

D.  $2/24 \approx 0.083$

**C — 0.125.** First increment the counts:  $N(s_{23}, \text{down}) \rightarrow 24$  and  $N(s_{23}, \text{down}, s_{24}) \rightarrow 3$ . Then  $P = 3/24 = 0.125$ .

# Active learning: explore vs exploit

Now the agent *chooses* its actions. Two competing pressures:

## Exploit

Take the action that maximizes the current  $V(s)$  (or  $Q(s, a)$ ). Reap known rewards.

## Explore

Try a different action. Gather data to learn a better model / better Q-values.

A purely greedy agent *seldom* converges to the optimal policy — the learned model is biased toward states the agent has already preferred.

# Exploration 1: $\epsilon$ -greedy



action probability


With probability  $\epsilon$ : take a **random** action (explore).

With probability  $1 - \epsilon$ : take the **greedy** action (exploit).

Start with a large  $\epsilon$  and **decay** it over time — explore early, exploit once the values settle.

Simple and popular, but it treats every sub-optimal action as equally worth trying.

# Exploration 2: softmax (Boltzmann)



Pick actions in proportion to their value, tuned by a **temperature**  $T$ :

$$P(a) = \frac{e^{Q(s,a)/T}}{\sum_{a'} e^{Q(s,a')/T}}$$

- Low  $T$ : nearly **greedy** (peaked on the best action).
- High  $T$ : nearly **uniform** (lots of exploration).

Unlike  $\epsilon$ -greedy, better actions are explored more than clearly-bad ones. Cool  $T$  over time (like simulated annealing).

# Exploration 3: optimistic initialization



Assume the best about actions you haven't tried enough. Replace a value  $u$  with an exploration function  $f$ :

$$f(u, n) = \begin{cases} R^+ & n < N_e \\ u & \text{otherwise} \end{cases}$$

$R^+$  is an optimistic reward bound;  $N_e$  is how many tries before we trust the real estimate  $u$ .

Every action looks great until tried  $N_e$  times, pulling the agent toward the unknown.

# Active ADP: bake exploration into the value

Active ADP = passive ADP, but now the agent *chooses* actions. The trick: act greedily on an **optimistic** value  $V^+$ , seen through the exploration function  $f$  from before.

Pick the action whose successor value — *viewed through*  $f$  — looks best:

$$a = \arg \max_a f \left( \underbrace{\sum_{s'} P(s' \mid s, a) V^+(s')}_{\text{estimated value of } a}, \underbrace{N(s, a)}_{\text{times tried}} \right)$$

$f$  inflates the value of **under-tried** actions ( $N(s, a) < N_e$ ), so the agent is pulled toward exploring them.

# The active ADP algorithm

## active-ADP( $N_e, R^+$ )

1. Initialize  $R(s)$ ,  $V^+(s)$ ,  $N(s, a)$ ,  $N(s, a, s')$ .
2. Repeat until every  $(s, a)$  is tried  $\geq N_e$  times and  $V^+$  converges:
  - Pick  $a = \arg \max_a f(\sum_{s'} P(s' | s, a) V^+(s'), N(s, a))$  (the rule above).
  - Execute  $a$ ; observe  $\langle s, a, s', r \rangle$ ; update  $R(s) \leftarrow r$ , counts, and  $P(s' | s, a)$ .
  - $V^+(s) \leftarrow R(s) + \gamma \max_a f(\sum_{s'} P(s' | s, a) V^+(s'), N(s, a))$ .
  - $s \leftarrow s'$ .

Same model-based loop as passive ADP — only the action choice and the **max** now pass through  $f$ .

# Recall: the Bellman equation for $V^*$

From L14, two connections between  $V^*$  and  $Q^*$ :

$$\textcircled{1} V^*(s) = \max_a Q^*(s, a)$$

$$\textcircled{2} Q^*(s, a) = R(s) + \gamma \sum_{s'} P(s' | s, a) V^*(s')$$

Substitute  $\textcircled{2}$  into  $\textcircled{1}$  (replace each  $Q^*$  with reward + future value):

$$V^*(s) = R(s) + \gamma \max_a \sum_{s'} P(s' | s, a) V^*(s')$$



# The Bellman equation for $Q^\star$

Now do the reverse — substitute ① ( $V^\star(s') = \max_{a'} Q^\star(s', a')$ ) into ②:

$$\begin{aligned} Q^\star(s, a) &= R(s) + \gamma \sum_{s'} P(s' \mid s, a) \underline{V^\star(s')} \\ &= R(s) + \gamma \sum_{s'} P(s' \mid s, a) \underline{\max_{a'} Q^\star(s', a')} \end{aligned}$$

Given  $Q^\star$ , the best action is  $\pi^\star(s) = \arg \max_a Q^\star(s, a)$  — no  $P$  needed. That is what lets  $Q$ -learning be **model-free**.

# Temporal difference (TD) error

After one transition  $\langle s, a, r, s' \rangle$ , pretend that transition was certain ( $P(s' | s, a) = 1$ ). Then by Bellman,  $Q(s, a)$  should equal the **target**:

$$Q(s, a) \stackrel{?}{=} \underbrace{R(s) + \gamma \max_{a'} Q(s', a')}_{\text{target}}$$

## TD error

$$\delta = \left( R(s) + \gamma \max_{a'} Q(s', a') \right) - Q(s, a)$$

How far the current estimate  $Q(s, a)$  is from the one-step target.

# From the TD error to the update

One sample is noisy, so don't jump all the way — nudge  $Q(s, a)$  toward the target by a fraction  $\alpha$  of the gap  $\delta$ :



$$\begin{aligned} Q(s, a) &\leftarrow Q(s, a) + \alpha \delta \\ &= Q(s, a) + \alpha \left( \underbrace{R(s) + \gamma \max_{a'} Q(s', a')}_{\text{target}} - Q(s, a) \right) \end{aligned}$$

$\alpha \rightarrow 0$ : barely move;  $\alpha = 1$ : jump straight to the target. This update rule is Q-learning.

# The Q-learning algorithm

## Q-learning( $\alpha, \gamma$ )

1. Initialize  $Q(s, a)$  arbitrarily; observe the starting state  $s$ .
2. Repeat until  $Q$  converges:
  - Pick  $a$  from  $s$  using *any* exploration strategy ( $\epsilon$ -greedy, softmax, optimistic).
  - Execute  $a$ ; observe reward  $r$  and next state  $s'$ .
  - $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a') - Q(s, a))$ .
  - $s \leftarrow s'$ .

**Model-free** — it never estimates  $P$ ; converges to  $Q^*$  if every  $(s, a)$  is visited infinitely often.

# Choosing the learning rate $\alpha$

The algorithm works for any fixed  $\alpha \in (0, 1]$ , but a good schedule lets  $\alpha$  **shrink** as a state–action pair is visited more often:

$$\alpha(N) = \frac{10}{9 + N(s, a)}$$

- Early visits (small  $N$ ):  $\alpha$  is large — move  $Q$  a lot on fresh information.
- Later visits (large  $N$ ):  $\alpha$  is small — fine-tune and average out noise.
- A decaying  $\alpha$  is exactly what guarantees convergence to  $Q^*$ .

# Q-learning: a single update

3x4 grid. Experience:  $\langle s_{32}, \text{right}, s_{33}, -0.04 \rangle$ . Take  $\gamma = 1, \alpha = 0.1$ .

|       | $s_{32}$ | $s_{33}$ |
|-------|----------|----------|
| up    | 0.55     | 0.4      |
| right | 0.7      | 0.9      |
| down  | 0.5      | 0.8      |
| left  | 0.4      | 0.5      |

Yellow cell = the one we update; green = the max-value action in  $s_{33}$ .

$$\text{target} = r + \gamma \cdot \max_{a'} Q(s_{33}, a') = -0.04 + 1.0 \cdot 0.9 = 0.86$$

$$\delta = \text{target} - Q(s_{32}, \text{right}) = 0.86 - 0.7 = 0.16$$

$$Q(s_{32}, \text{right}) \leftarrow 0.7 + 0.1 \cdot 0.16 = 0.716$$

# SARSA: on-policy temporal difference

Observe a 5-tuple  $\langle s, a, r, s', a' \rangle$  where  $a'$  is the action **actually taken** in  $s'$  (per the current policy):

## Q-learning (off-policy)

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

Target uses the *greedy* next action. Learns about the optimal policy regardless of what's actually played.

## SARSA (on-policy)

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$$

Target uses the *actually-played* next action  $a'$ . Learns about the policy the agent is following (exploration and all).

SARSA = State, Action, Reward, State, Action.

# Why is Q-learning *off-policy*?

Every step, the agent runs **two** policies at once:

- **Behavior policy** — how it *actually* picks actions ( $\epsilon$ -greedy: usually greedy, sometimes explores).
- **Target policy** — the policy whose value the update is learning about.

## Q-learning

The update uses  $\max_{a'} Q(s', a')$  — the value of the **greedy** next action — even when the agent is about to *explore* a different one.

Behavior  $\neq$  target  $\Rightarrow$  **off-policy**. Converges to  $Q^*$  no matter how it explores.

## SARSA

The update uses  $Q(s', a')$  for the action **actually chosen** by  $\epsilon$ -greedy — the very move it will play next.

Behavior = target  $\Rightarrow$  **on-policy**. Learns the value of the exploring policy itself.

**Off-policy:** learn about one policy (greedy) while *behaving* by another (exploratory). **On-policy:** learn about the policy you actually run.



# SARSA vs Q-learning: the scenario

Update  $Q(s_{13}, \text{down}) = 0.4$  after  $\langle s_{13}, \text{down}, s_{23} \rangle$ , with  $r = -0.04$ ,  $\gamma = 0.9$ ,  $\alpha = 0.4$ .

The  $\varepsilon$ -greedy step explores  $a' = \text{right}$  in  $s_{23}$  (into the bad cell).

| $Q(s_{23}, \cdot)$ | up  | right | down | left |
|--------------------|-----|-------|------|------|
|                    | 0.3 | -0.4  | 0.7  | 0.3  |

Yellow = the explored next action  $a' = \text{right}$  (SARSA's target); green =  $\max_{a'} Q$  (Q-learning's target).

# SARSA vs Q-learning: the two updates

**SARSA: uses  $a' = \text{right}$**

$$\text{target} = -0.04 + 0.9 \cdot (-0.4) = -0.4$$

$$\delta = -0.4 - 0.4 = -0.8$$

$$Q(s_{13}, \text{down}) \leftarrow 0.4 + 0.4 \cdot (-0.8)$$

**= 0.08 (punished for the bad step)**

**Q-learning: uses  $\max_{a'} = 0.7$**

$$\text{target} = -0.04 + 0.9 \cdot 0.7 = 0.59$$

$$\delta = 0.59 - 0.4 = 0.19$$

$$Q(s_{13}, \text{down}) \leftarrow 0.4 + 0.4 \cdot 0.19$$

**= 0.476 (rewarded for the best)**

The agent is about to play  $a' = \text{right}$  ( $-0.4$ ), but Q-learning updates toward  $\max = 0.7$  (down) — not the move it will make. That mismatch is "off-policy."

SARSA is **conservative** (learns what really happens); Q-learning is **aggressive** (learns the optimal policy).

# When to use what (Q2)

**Q2.** For a high-stakes *online* learning problem (for example, a self-driving car still in training), which algorithm is safer?

- A. Q-learning
- B. **SARSA**
- C. I don't know

**B — SARSA.** SARSA's target uses the action that will *actually* be taken while exploring, so it learns to avoid catastrophic outcomes along the exploration path. Q-learning's target uses the greedy action, ignoring the cost of exploration mistakes.

# Putting the labels together (Q3)

**Q3.** For each of **Q-learning** and **SARSA**: is it *model-based* or *model-free*? Is it *on-policy* or *off-policy*?

**Both are model-free** (neither estimates  $P$ ). **Q-learning is off-policy** — its target uses  $\max_{a'} Q(s', a')$ , the greedy action. **SARSA is on-policy** — its target uses the action actually played.

# ADP vs Q-learning vs SARSA

|                            | ADP                    | Q-learning              | SARSA                   |
|----------------------------|------------------------|-------------------------|-------------------------|
| Learns transition $P$ ?    | yes (model-based)      | no (model-free)         | no (model-free)         |
| On / off policy            | $n/a$ (eval-only)      | off-policy              | on-policy               |
| Computation per experience | heavy (solve Bellman)  | cheap (one TD update)   | cheap (one TD update)   |
| Convergence speed          | fast (few experiences) | slower, higher variance | slower, higher variance |

ADP is sample-efficient but expensive per step; TD methods (Q, SARSA) flip that trade-off.

## Learning goals (recap)

- ✓ Trace + implement **passive ADP**.
- ✓ Explain the **exploration vs exploitation** trade-off.
- ✓ Trace + implement **active ADP**.
- ✓ Trace + implement **Q-learning**.
- ✓ Trace + implement **SARSA**.
- ✓ Tell **on-policy** from **off-policy**.

Search ›

Uncertainty ›

Decisions ›

Learning

## Next module: Machine Learning and Deep Learning

RL was learning from rewards. The last module turns to *learning from data*.

- **L16** — supervised learning, bias / variance, cross-validation.
- **L17** — unsupervised: k-means, PCA, autoencoders, GANs.
- **L18** — decision trees: entropy, information gain, ID3.
- **L19–L21** — neural networks: forward / backward / optimizers.