

# CS 486/686

# Reinforcement Learning

Yuntian Deng

---

Lecture 15

RN 22.1–22.3 · PM 12.1, 12.5, 12.8

Search ›

Uncertainty ›


**Decisions** ›

Learning

## Learning goals

- Trace + implement **passive adaptive dynamic programming** (passive ADP).
- Explain the **exploration vs exploitation** trade-off.
- Trace + implement **active ADP**.
- Trace + implement **Q-learning**.
- Trace + implement **SARSA**.
- Tell **on-policy** from **off-policy** learning.

# The RL setting: no $P$ , no $R$ up front



An MDP whose **transition  $P$**  and **reward  $R$**  are unknown. The agent sees state  $s$  and reward  $r$ , picks an action  $a$ , and observes  $s'$ . Goal: maximize discounted return.

## Three new challenges:

- **Credit assignment:** which past action caused this reward?
- **Long-term impact:** how will this action affect future returns?
- **Explore vs exploit:** try new actions or stick with the best known?

# Passive learning: just evaluate a fixed policy

Fix a policy  $\pi$ . Goal: learn  $V^\pi(s)$  for every state  $s$ .

## Adaptive Dynamic Programming (ADP)

1. **Learn** the transition  $P(s' \mid s, a)$  and reward  $R(s)$  from observed experience.
2. **Solve** the Bellman policy-evaluation equations:

$$V^\pi(s) = R(s) + \gamma \sum_{s'} P(s' \mid s, \pi(s)) V^\pi(s').$$

A **model-based** approach — it builds an explicit model of the world.

# The passive ADP algorithm

## passive-ADP( $\pi$ )

1. Initialize  $R(s)$ ,  $N(s, a)$ ,  $N(s, a, s')$ ,  $V^\pi(s)$ .
2. Follow  $\pi$  and observe an experience  $\langle s, a, s', r \rangle$ .
3. Update reward:  $R(s) \leftarrow r$ .
4. Update counts and transition probability:

$$N(s, a) += 1, \quad N(s, a, s') += 1, \quad P(s' | s, a) = \frac{N(s, a, s')}{N(s, a)}.$$

5. Re-solve Bellman to get  $V^\pi(s)$  (exactly or iteratively).

Each new experience refines the model; each Bellman solve costs  $O(n^3)$  or a few iterations.

# Passive ADP: a worked example

|          |    |
|----------|----|
| $s_{11}$ | +1 |
| $s_{21}$ | -1 |

$\pi(s_{11}) = \text{down}, \pi(s_{21}) = \text{right}, \gamma = 0.9.$

Before:  $N(s, a) = 5$  for all  $s, a$ ;  $N(s, a, s') = 3$  for intended direction, 1 for each side.

New experience:  $\langle s_{11}, \text{down}, s_{21}, -0.04 \rangle$ .

Update counts:  $N(s_{11}, \text{down}) = 6, N(s_{11}, \text{down}, s_{21}) = 4.$

Solve Bellman with the refreshed model:

$$V(s_{11}) = -0.4378, \quad V(s_{21}) = -0.8034.$$

## Quick check: probability update

**Q1.** In a 3x4 grid, suppose  $N(s_{23}, \text{down}) = 23$  and  $N(s_{23}, \text{down}, s_{24}) = 2$ . The agent tries to move *down* but slips and lands in  $s_{24}$ . What is the updated  $P(s_{24} \mid s_{23}, \text{down})$ ?

- A.  $2/23 \approx 0.087$
- B.  $3/23 \approx 0.130$
- C.  $3/24 = 0.125$
- D.  $2/24 \approx 0.083$

**C — 0.125.** First increment the counts:  $N(s_{23}, \text{down}) \rightarrow 24$  and  $N(s_{23}, \text{down}, s_{24}) \rightarrow 3$ . Then  $P = 3/24 = 0.125$ .

# Active learning: explore vs exploit

Now the agent *chooses* its actions. Two competing pressures:

## Exploit

Take the action that maximizes the current  $V(s)$  (or  $Q(s, a)$ ). Reap known rewards.

## Explore


Try a different action. Gather data to learn a better model / better Q-values.

A purely greedy agent *seldom* converges to the optimal policy — the learned model is biased toward states the agent has already preferred.



# Three exploration strategies

## $\epsilon$ -greedy




Random with prob  $\epsilon$ ;  
greedy with prob  $1 - \epsilon$ .

Decay  $\epsilon$  over time.

Simple, but treats all sub-optimal actions equally.


## Softmax / Boltzmann



$$P(a) = \frac{e^{Q(s,a)/T}}{\sum_{a'} e^{Q(s,a')/T}}$$

Low  $T$ : peaky / greedy. High  $T$ :  
uniform / exploratory.

## Optimistic init



$$f(u, n) = \begin{cases} R^+ & n < N_e \\ u & \text{else} \end{cases}$$

"Shiny new toy" bonus until  
each action has been tried  $N_e$   
times.

# The active ADP algorithm

## **active-ADP( $N_e, R^+$ )**

1. Initialize  $R(s)$ ,  $V^+(s)$ ,  $N(s, a)$ ,  $N(s, a, s')$ .
2. Repeat until every  $(s, a)$  is visited  $\geq N_e$  times and  $V^+$  converged:
  - Pick action  $a = \arg \max_a f(\sum_{s'} P(s' | s, a) V^+(s'), N(s, a))$ .
  - Execute  $a$ ; observe  $\langle s, a, s', r \rangle$ ; update  $R(s) \leftarrow r$ , counts, and  $P(s' | s, a)$ .
  - Update  $V^+(s) \leftarrow R(s) + \gamma \max_a f(\sum_{s'} P(s' | s, a) V^+(s'), N(s, a))$ .

Same model-based loop as passive ADP, plus the exploration bonus  $f$  baked into action selection.

# From V to Q: learn without a model

Bellman for  $V$  and  $Q$  are both fixed-point equations:

## Bellman for $V^\star$

$$V^\star(s) = R(s) + \gamma \max_a \sum_{s'} P(s' | s, a) V^\star(s')$$

To pick  $\pi^\star$  you still need to know  $P$  and combine it with  $V^\star$ .

## Bellman for $Q^\star$

$$Q^\star(s, a) = R(s) + \gamma \sum_{s'} P(s' | s, a) \max_{a'} Q^\star(s', a')$$

If we know  $Q^\star$ , then  $\pi^\star(s) = \arg \max_a Q^\star(s, a)$  — no  $P$  needed.

**model-free!**

Learning  $Q$  directly is the trick that makes RL practical at scale.

# Temporal difference (TD) error

After one transition  $\langle s, a, r, s' \rangle$ , pretend that transition was certain ( $P(s' | s, a) = 1$ ). Then by Bellman:

$$Q(s, a) \stackrel{?}{=} \underbrace{R(s) + \gamma \max_{a'} Q(s', a')}_{\text{target}}$$

## TD error

$$\delta = \left( R(s) + \gamma \max_{a'} Q(s', a') \right) - Q(s, a)$$



Idea: nudge  $Q(s, a)$  toward the target by an amount proportional to  $\delta$ .

# The Q-learning algorithm

## Q-learning( $\alpha, \gamma$ )

1. Initialize  $Q(s, a)$  arbitrarily,  $R(s)$ ,  $N(s, a)$ .
2. Repeat until  $Q$  converges:
  - Pick  $a$  from  $s$  using *any* exploration strategy ( $\epsilon$ -greedy, softmax, optimistic).
  - Execute  $a$ ; observe  $\langle s, a, s', r \rangle$ ; update  $R(s) \leftarrow r$ .
  - $Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a') - Q(s, a))$ .
  - $s \leftarrow s'$ .

- **Model-free:** never estimates  $P$  explicitly.
- **Learning rate:** smaller  $\alpha$  = slower but more accurate; e.g.  $\alpha(N) = 10/(9 + N(s, a))$ .
- Converges to  $Q^*$  if every  $(s, a)$  is visited infinitely often.

# Q-learning: a single update

3x4 grid. Experience:  $\langle s_{32}, \text{right}, s_{33}, -0.04 \rangle$ . Take  $\gamma = 1, \alpha = 0.1$ .

|       | $s_{32}$ | $s_{33}$ |
|-------|----------|----------|
| up    | 0.55     | 0.4      |
| right | 0.7      | 0.9      |
| down  | 0.5      | 0.8      |
| left  | 0.4      | 0.5      |

Yellow cell = the one we update; green = the max-value action in  $s_{33}$ .

$$\text{target} = r + \gamma \cdot \max_{a'} Q(s_{33}, a') = -0.04 + 1.0 \cdot 0.9 = 0.86$$

$$\delta = \text{target} - Q(s_{32}, \text{right}) = 0.86 - 0.7 = 0.16$$

$$Q(s_{32}, \text{right}) \leftarrow 0.7 + 0.1 \cdot 0.16 = 0.716$$

# SARSA: on-policy temporal difference

Observe a 5-tuple  $\langle s, a, r, s', a' \rangle$  where  $a'$  is the action **actually taken** in  $s'$  (per the current policy):

## Q-learning (off-policy)

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

Target uses the *greedy* next action. Learns about the optimal policy regardless of what's actually played.

## SARSA (on-policy)

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$$

Target uses the *actually-played* next action  $a'$ . Learns about the policy the agent is following (exploration and all).

SARSA = State, Action, Reward, State, Action.

# SARSA vs Q-learning: a side-by-side update

At  $s_{13}$ , update  $Q(s_{13}, \text{down}) = 0.4$ . Next state  $s_{23}$ :  $Q = (\text{L}:0.3, \text{R}:-0.4, \text{D}:0.7, \text{U}:0.3)$ ,  $r = -0.04$ ,  $\gamma = 0.9$ ,  $\alpha = 0.4$ . The agent's  $\epsilon$ -greedy step explores  $a' = \text{right}$  (into the bad cell).

**SARSA: uses  $a' = \text{right}$**

$$\begin{aligned}\text{target} &= -0.04 + 0.9 \cdot (-0.4) = -0.4 \\ \delta &= -0.4 - 0.4 = -0.8 \\ Q(s_{13}, \text{down}) &\leftarrow 0.4 + 0.4 \cdot (-0.8)\end{aligned}$$

**= 0.08 (punished for the bad step)**

**Q-learning: uses  $\max_{a'} = 0.7$**

$$\begin{aligned}\text{target} &= -0.04 + 0.9 \cdot 0.7 = 0.59 \\ \delta &= 0.59 - 0.4 = 0.19 \\ Q(s_{13}, \text{down}) &\leftarrow 0.4 + 0.4 \cdot 0.19\end{aligned}$$

**= 0.476 (rewarded for the best)**

SARSA is **conservative** (learns what really happens); Q-learning is **aggressive** (learns the optimal policy).



# ADP vs Q-learning vs SARSA

|                            | ADP                    | Q-learning              | SARSA                   |
|----------------------------|------------------------|-------------------------|-------------------------|
| Learns transition $P$ ?    | yes (model-based)      | no (model-free)         | no (model-free)         |
| On / off policy            | $n/a$ (eval-only)      | off-policy              | on-policy               |
| Computation per experience | heavy (solve Bellman)  | cheap (one TD update)   | cheap (one TD update)   |
| Convergence speed          | fast (few experiences) | slower, higher variance | slower, higher variance |

ADP is sample-efficient but expensive per step; TD methods (Q, SARSA) flip that trade-off.

# When to use what (Q2)

**Q2.** For a high-stakes *online* learning problem (for example, a self-driving car still in training), which algorithm is safer?

- A. Q-learning
- B. **SARSA**
- C. I don't know

**B — SARSA.** SARSA's target uses the action that will *actually* be taken while exploring, so it learns to avoid catastrophic outcomes along the exploration path. Q-learning's target uses the greedy action, ignoring the cost of exploration mistakes.

# Putting the labels together (Q3)

Q3. Which statement is true?

- A. Q-learning is model-free + off-policy; SARSA is model-based + on-policy.
- B. Q-learning is model-based + off-policy; SARSA is model-free + on-policy.
- C. Q-learning is model-free + on-policy; SARSA is model-based + off-policy.
- D. Q-learning is model-based + on-policy; SARSA is model-free + off-policy.
- E. **None of the above.**

**E — None of the above.** Both Q-learning and SARSA are *model-free*. Q-learning is off-policy (target uses  $\max_{a'} Q$ ); SARSA is on-policy (target uses the action actually played).

## Learning goals (recap)

- ✓ Trace + implement **passive ADP**.
- ✓ Explain the **exploration vs exploitation** trade-off.
- ✓ Trace + implement **active ADP**.
- ✓ Trace + implement **Q-learning**.
- ✓ Trace + implement **SARSA**.
- ✓ Tell **on-policy** from **off-policy**.

Search ›

Uncertainty ›

Decisions ›

Learning

## Next module: Machine Learning and Deep Learning

RL was learning from rewards. The last module turns to *learning from data*.

- **L16** — supervised learning, bias / variance, cross-validation.
- **L17** — unsupervised: k-means, PCA, autoencoders, GANs.
- **L18** — decision trees: entropy, information gain, ID3.
- **L19–L21** — neural networks: forward / backward / optimizers.