CS 486/686
Training Models From Scratch

Yuntian Deng

Lecture 17

Core loop: data → model → loss → gradient → update

Recorded lecture (asynchronous)

> 🎥 This lecture is **pre-recorded** - I'm away at **ICML** this week, so there's no in-person class today. Watch at your own pace.
>
> 📅 **Due today (Tue Jul 7):** Chat 8. The CS686 project proposal is now due **Thu Jul 9** (submit on LEARN under "Project Proposal").
>
> 💬 Questions? Post on **Piazza** - the TAs are available, and I'll follow up when I'm back.

Today's target

By the end, this loop should feel readable:

```
for x, y in loader:
    pred = model(x)
    loss = loss_fn(pred, y)
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()
```

Everything later - neural nets, LLMs, diffusion - scales up this idea.

The learning recipe

Parameters are the knobs

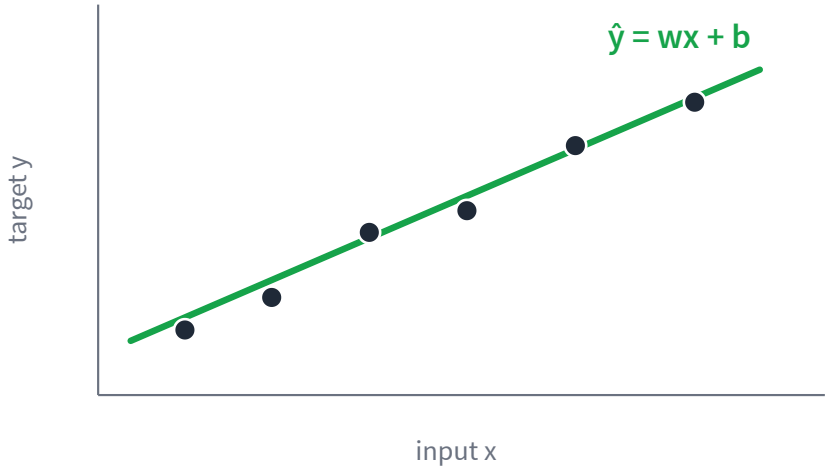

The data is fixed.

The model has tunable **parameters**: w and b.

Training means choosing parameters that make predictions good.

Start with one knob

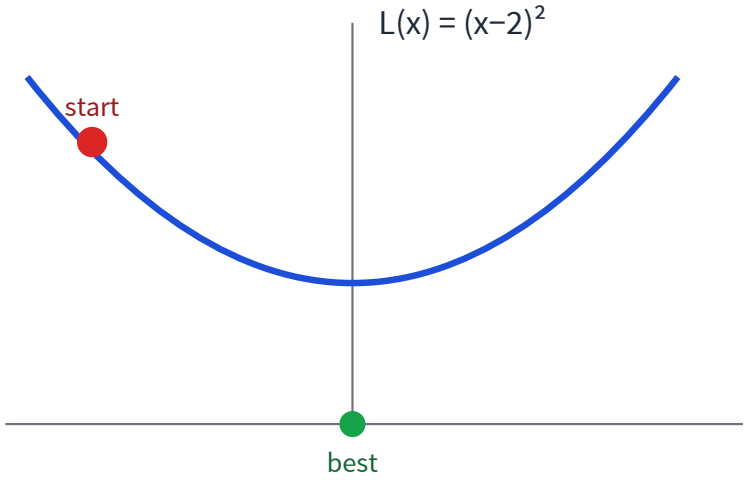

Forget ML for one slide.

We only need to make a loss small.

This is the smallest possible version of training.

Move downhill

The gradient points uphill, so step the other way.

$$x \leftarrow x - \eta \frac{dL}{dx}$$

$-dL/dx$

direction that decreases loss

η

learning rate: how big a step?

Playground: gradient descent LIVE

Interactive demo — runs live in your browser (open the slides to try it).

Interactive on the live slides: drag the learning rate η and press Step / Run to watch x roll down $L(x) = (x - 2)^2$. Small steps crawl; too large and it diverges. Runs instantly in your browser.

PyTorch computes gradients for us

Manual derivative

```
def loss(x):
    return (x - 2) ** 2

def manual_grad(x):
    return 2 * (x - 2)
```

Automatic differentiation

```
x = nn.Parameter(torch.tensor(-4.0))
L = (x - 2) ** 2
L.backward()
print(x.grad)  # -12
```

For millions of parameters, we rely on autograd. But how does it know the answer?

How can gradients be automatic?

Every operation records its own **local derivative**; the **chain rule** multiplies them backward.

Forward: $u = x - 2$, then $L = u^2$.

Chain rule: $\frac{dL}{dx} = \frac{dL}{du} \cdot \frac{du}{dx}$

Local: $\frac{dL}{du} = 2u$ (since $L = u^2$)

Local: $\frac{du}{dx} = 1$ (since $u = x - 2$)

Multiply: $\frac{dL}{dx} = 2u \cdot 1 = 2(x - 2)$

At $x = -4$: $\frac{dL}{dx} = 2(-6) = -12$ — exactly `x.grad`.

Autograd follows the computation graph

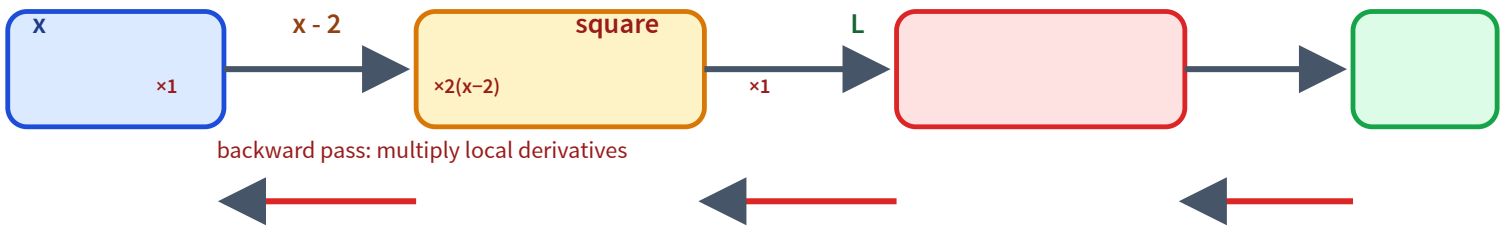

Run it: autograd from scratch LIVE PYTHON

Interactive demo — runs live in your browser (open the slides to try it).

Interactive on the live slides: a ~20-line reverse-mode autograd (a tiny `Value` class) builds the graph $x \to (x - 2) \to \text{square} \to L$, runs `L.backward()`, and prints `x.grad == -12`. Editable real Python, executed in your browser via Pyodide.

From one knob to many

A model has many parameters, collected as θ. The loop does not change.

$$\theta \leftarrow \theta - \eta \nabla_\theta L(\theta)$$

```
optimizer = SGD(model.parameters(), lr=0.1)

for x, y in loader:
    pred = model(x)
    loss = loss_fn(pred, y)
    loss.backward()
    optimizer.step()
    optimizer.zero_grad()
```

`model.parameters()` → the parameters θ

`loss.backward()` → the gradient $\nabla_\theta L$

`optimizer.step()` → update $\theta \leftarrow \theta - \eta \nabla_\theta L$

Example 1: regression

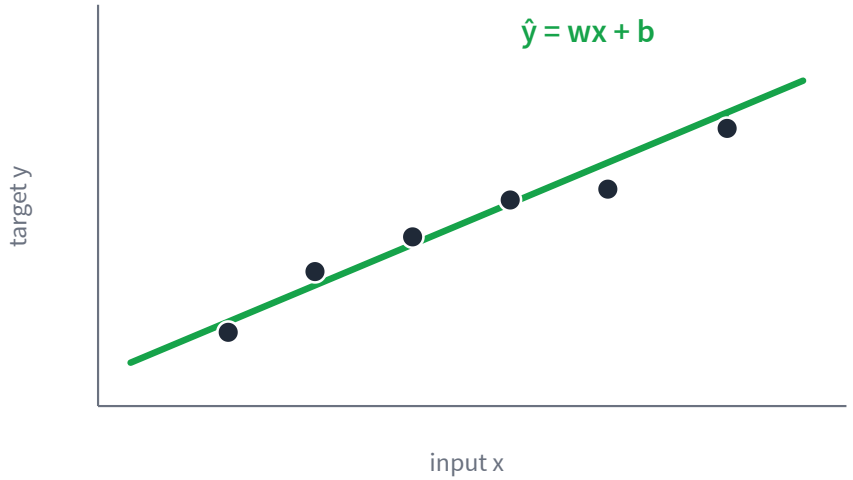

Prediction: $\hat{y} = wx + b$ (w, b are the parameters)

Error: $L = \frac{1}{m} \sum_i (\hat{y}_i - y_i)^2$

Start with a bad line — big errors (dashed).

Measure the error, nudge w, b to shrink it.

Repeat — the errors keep shrinking.

Good fit: the errors are small.

Regression in PyTorch

```
# the knob: a line y = wx + b
model = nn.Linear(1, 1)
# how we score error: mean squared error
loss_fn = nn.MSELoss()
# how we update the parameters
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
# forward: predict from inputs
pred = model(x)
# measure how wrong we are
loss = loss_fn(pred, y)
# autograd fills in the gradients
loss.backward()
# one step downhill: update w and b
optimizer.step()
# clear gradients before the next batch
optimizer.zero_grad()
```

Run it: train the line LIVE PYTHON

Interactive demo — runs live in your browser (open the slides to try it).

Interactive on the live slides: a real numpy training loop (predict $\longrightarrow$ MSE $\longrightarrow$ gradients $\longrightarrow$ update) fits $\hat{y} = wx + b$ to noisy data. Drag the learning rate and steps, press Train, and watch the line snap to the data as the loss drops. Runs in your browser via Pyodide.

Example 2: classification

Now the target is a **category**, not a number.

Regression

$\hat{y} = 217.3$

a number

Classification

$\hat{y} =$ “spam”

a class

But gradient descent needs a smooth output — so we predict a **probability** instead.

Logistic regression: predict a probability

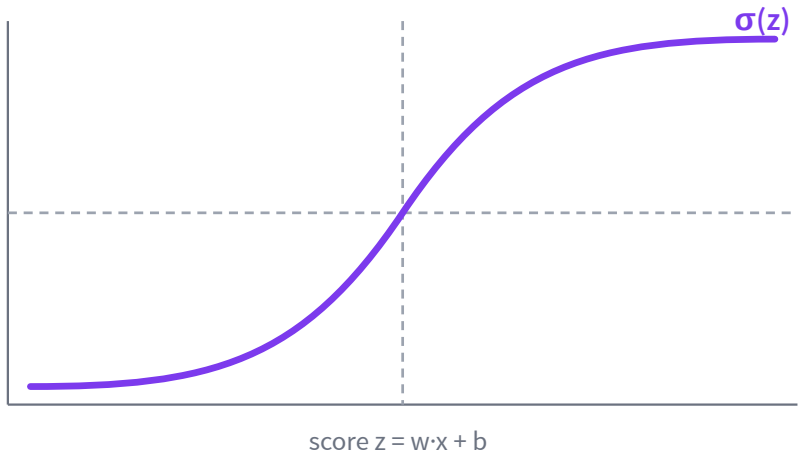

Compute a score:

$$z = w \cdot x + b$$

Squash it into a probability:

$$P(\text{spam} \mid x) = \sigma(z) = \frac{1}{1+e^{-z}}$$

Loss for classification

Cross-entropy punishes being confidently wrong.

Recall $\hat{y} = P(y = 1 \mid x) = \sigma(z)$ — the probability the model predicts.

$$L = -\left[y \log \hat{y} + (1 - y) \log(1 - \hat{y})\right]$$

True label spam $\Rightarrow y = 1$, so $L = -\log \hat{y}$.

confident & right

$\hat{y} = 0.9$

$L = -\log 0.9 \approx 0.11$

confident & wrong

$\hat{y} = 0.1$

$L = -\log 0.1 \approx 2.30$

Playground: cross-entropy loss LIVE

Interactive demo — runs live in your browser (open the slides to try it).

Interactive on the live slides: drag the score z to see the sigmoid probability $\sigma(z)$ and the cross-entropy loss — being confident and wrong is punished hard. Flip the true label to compare. Runs instantly in your browser.

Many classes: softmax

One score per class; softmax turns scores into probabilities.

$$P(y = k \mid x) = \frac{e^{z_k}}{\sum_j e^{z_j}}$$

Next-token prediction is exactly this, over a vocabulary.

Playground: softmax & temperature LIVE

Interactive demo — runs live in your browser (open the slides to try it).

Interactive on the live slides: adjust the class scores (logits) and a temperature slider and watch the softmax probabilities sharpen or flatten. Runs instantly in your browser.

The full loop, annotated

```
# loop over the training data
for x, y in loader:
    # forward pass: predict
    pred = model(x)
    # measure the error
    loss = loss_fn(pred, y)
    # backward pass: compute gradients
    loss.backward()
    # update the parameters
    optimizer.step()
    # reset gradients for the next step
    optimizer.zero_grad()
```

This is the pattern to recognize for the rest of the module.

Why mini-batches?

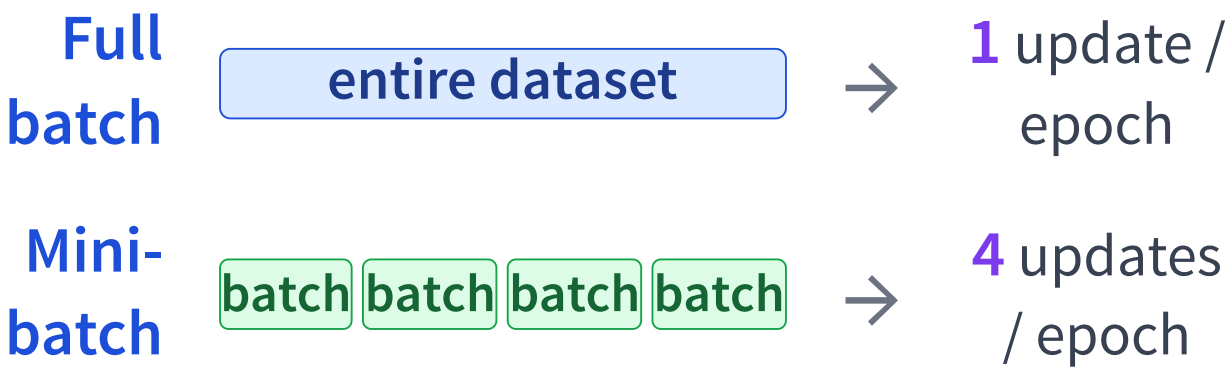

Full batch: few, expensive, smooth steps. Mini-batch: many, cheap, noisy steps — usually faster progress, and GPU-friendly.

Split the data: train / validation / test

train

updates the weights

validation

chooses settings

test

final, one-time estimate

Never tune on the test set.

More power is not always better

Same data, increasingly flexible models.

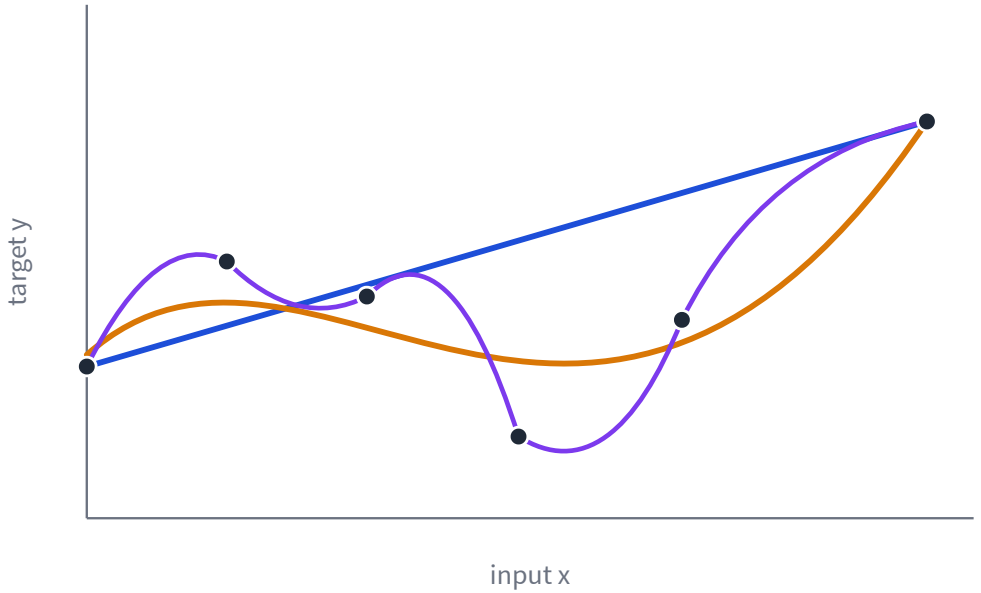

Watch train and validation loss

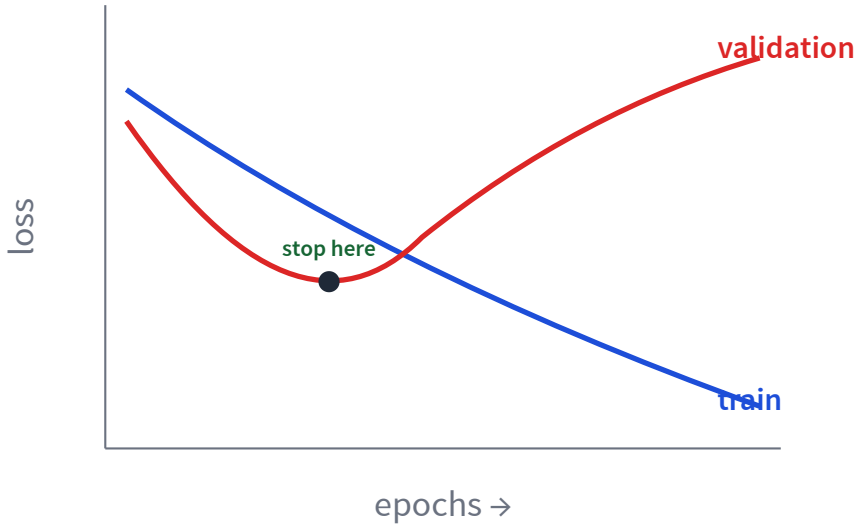

Training loss keeps dropping.

Validation loss starts **rising** — that growing gap is overfitting. Stop where validation bottoms out.

Playground: overfitting LIVE PYTHON

Interactive demo — runs live in your browser (open the slides to try it).

Interactive on the live slides: raise the polynomial degree and watch training loss keep dropping while validation loss turns up - overfitting appears as a U-shaped validation curve. Fit with numpy in your browser via Pyodide.

Three ways to fight overfitting

More data

show the model more of the world

Regularize

prefer simpler weights

$$L(\theta) + \lambda \|\theta\|^2$$

Early stop

stop when validation turns up

What carries forward?

Only the **model** and the **loss** change. The loop stays the same.

Task	Model (architecture)	Loss
Regression (today)	Linear / MLP	squared error
Language model	Transformer	cross-entropy (next token)
Image generation	Diffusion U-Net	squared error (predict noise)

forward → loss → backward → step, every time.

Recap - next: learned representations

✓ A model learns by minimizing a **loss**.

✓ **Autograd** computes gradients through a graph.

✓ Regression uses squared error; classification uses **cross-entropy**.

✓ Validation loss catches **overfitting**.

L18: if linear models use fixed features, how do neural nets learn their own?