

CS 486/686

Training Models From Scratch

Yuntian Deng

Lecture 17

Core loop: data $\rightarrow$ model $\rightarrow$ loss $\rightarrow$ gradient $\rightarrow$ update

Recorded lecture (asynchronous)



This lecture is **pre-recorded** - I'm away at **ICML** this week, so there's no in-person class today. Watch at your own pace.



Due today (Tue Jul 7): Chat 8. The CS686 project proposal is now due **Thu Jul 9** (submit on LEARN under "Project Proposal").



Questions? Post on **Piazza** - the TAs are available, and I'll follow up when I'm back.

Today's target

By the end, this loop should feel readable:


```
for x, y in loader:  
    pred = model(x)  
    loss = loss_fn(pred, y)  
    loss.backward()  
    optimizer.step()  
    optimizer.zero_grad()
```

Everything later - neural nets, LLMs, diffusion - scales up this idea.

The learning recipe



Parameters are the knobs




The data is fixed.

The model has tunable **parameters**: w and b .

Training means choosing parameters that make predictions good.

Start with one knob



Forget ML for one slide.

We only need to make a loss small.

This is the smallest possible version of training.

Move downhill

The gradient points uphill, so step the other way.

$$x \leftarrow x - \eta \frac{dL}{dx}$$

$$-dL/dx$$

direction that decreases loss

$$\eta$$

learning rate: how big a step?

PyTorch computes gradients for us

Manual derivative

```
def loss(x):  
    return (x - 2) ** 2  
  
def manual_grad(x):  
    return 2 * (x - 2)
```

Automatic differentiation

```
x = nn.Parameter(torch.tensor(-4.0))  
L = (x - 2) ** 2  
L.backward()  
print(x.grad)  # -12
```

For millions of parameters, we rely on autograd. But how does it know the answer?

How can gradients be automatic?

Every operation is recorded, and each knows its own **local derivative**. The **chain rule** multiplies them backward.

Forward

$$u = x - 2$$

$$L = u^2$$


Backward (chain rule)

$$\frac{dL}{du} = 2u, \quad \frac{du}{dx} = 1$$

$$\frac{dL}{dx} = 2u \cdot 1 = 2(x - 2)$$

At $x = -4$: $\frac{dL}{dx} = 2(-6) = -12$ — exactly x . grad.

Autograd follows the computation graph



From one knob to many


A model has many parameters, collected as θ . The loop does not change.

```
for x, y in loader:
    pred = model(x)
    loss = loss_fn(pred, y)
    loss.backward()      #  $\nabla_{\theta} L$ 
    optimizer.step()     #  $\theta \leftarrow \theta - \eta \nabla_{\theta} L$ 
    optimizer.zero_grad()
```

$$\theta \leftarrow \theta - \eta \nabla_{\theta} L(\theta)$$

`model.parameters()` $\rightarrow \theta$
`loss.backward()` $\rightarrow \nabla_{\theta} L$
`optimizer.step()` $\rightarrow$ update θ

Example 1: regression



Prediction is a number:

$$\hat{y} = w x + b$$

w is the slope, b the intercept — the knobs we learn.

$$L = \frac{1}{m} \sum_i (\hat{y}_i - y_i)^2$$

Squared error punishes large residuals (dashed).

Regression in PyTorch

```
# the knob: a line  $y = wx + b$ 
```

```
model = nn.Linear(1, 1)
```

```
# how we score error: mean squared error
```

```
loss_fn = nn.MSELoss()
```

```
# how we update the knobs
```

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
```

```
# forward: predict from inputs
```

```
pred = model(x)
```

```
# measure how wrong we are
```

```
loss = loss_fn(pred, y)
```

```
# autograd fills in the gradients
```

```
loss.backward()
```

```
# one step downhill: update w and b
```

```
optimizer.step()
```

```
# clear gradients before the next batch
```

```
optimizer.zero_grad()
```

Example 2: classification

Now the target is a **category**, not a number.

Regression


$\hat{y} = 217.3$
a number

Classification

$\hat{y} = \text{"spam"}$
a class

But gradient descent needs a smooth output — so we predict a **probability** instead.

Logistic regression: predict a probability



Compute a score:

$$z = w^{\top} x + b$$

Squash it into a probability:

$$P(\text{spam} \mid x) = \sigma(z) = \frac{1}{1 + e^{-z}}$$

Loss for classification

Cross-entropy punishes being confidently wrong.

$$L = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

True label spam $\Rightarrow y = 1$, so $L = -\log \hat{y}$.

confident & right

$$\hat{y} = 0.9$$

$$L = -\log 0.9 \approx 0.11$$

confident & wrong

$$\hat{y} = 0.1$$

$$L = -\log 0.1 \approx 2.30$$

Many classes: softmax

One score per class; softmax turns scores into probabilities.

$$P(y = k \mid x) = \frac{e^{z_k}}{\sum_j e^{z_j}}$$




Next-token prediction is exactly this, over a vocabulary.

The full loop, annotated

```
# grab the next mini-batch
for x, y in loader:
# forward pass: predict
    pred = model(x)
# measure the error
    loss = loss_fn(pred, y)
# backward pass: compute gradients
    loss.backward()
# update the parameters
    optimizer.step()
# reset gradients for the next step
    optimizer.zero_grad()
```

This is the pattern to recognize for the rest of the module.

Why mini-batches?



Full batch: few, expensive, smooth steps. Mini-batch: many, cheap, noisy steps — usually faster progress, and GPU-friendly.

Split the data: train / validation / test

train

updates the weights

validation

chooses settings


test

final, one-time
estimate

Never tune on the test set.


More power is not always better

Same data, increasingly flexible models.



— underfit — good fit — overfit

Watch train and validation loss



Training loss keeps dropping.

Validation loss starts **rising** — that growing gap is overfitting. Stop where validation bottoms out.

Three ways to fight overfitting

More data

show the model
more of the world

Regularize

prefer simpler
weights

$$L(\theta) + \lambda \|\theta\|^2$$

Early stop

stop when validation
turns up

What carries forward?

Only the **model** and the **loss** change. The loop stays the same.

Task	Model (architecture)	Loss
Regression (today)	Linear / MLP	squared error
Language model	Transformer	cross-entropy (next token)
Image generation	Diffusion U-Net	squared error (predict noise)

forward → loss → backward → step, every time.

Recap - next: learned representations

- ✓ A model learns by minimizing a **loss**.
- ✓ **Autograd** computes gradients through a graph.
- ✓ Regression uses squared error; classification uses **cross-entropy**.
- ✓ Validation loss catches **overfitting**.

L18: if linear models use fixed features, how do neural nets learn their own?