

# CS 486/686

# Neural Networks and Learned Embeddings

Yuntian Deng

---

Lecture 18

From fixed features to learned vector representations

# Recorded lecture (asynchronous)



This lecture is **pre-recorded** - I'm away at **ICML** this week, so there's no in-person class today. Watch at your own pace.



**Due today (Thu Jul 9):** Assignment 2. Chat 9 is now out.



Questions? Post on **Piazza** - the TAs are available, and I'll follow up when I'm back.

# The problem with fixed features

L17 assumed the useful input vector already exists.

**pixels**

Where is the cat ear?

**word IDs**


Why is dog close to  
puppy?

**screenshots**

Which region is a  
button?


Modern ML learns the features too.

# Neural nets learn representations



The hidden layer is not just computation: it is a learned view of the input.

# A neuron is a feature detector



$$z = w^{\top} x + b$$

$$a = g(z)$$

Weights decide what pattern to detect.

# Why add a nonlinearity?

Stacking linear layers is still linear.

$$W_2(W_1x) = (W_2W_1)x$$

**ReLU**


$$g(x) = \max(0, x)$$

**GELU / SiLU**

smooth modern defaults

Nonlinearity lets layers bend the representation.

# One hidden layer = learned feature map




Each hidden unit learns a detector.

The output layer combines detectors.

This is how fixed features become learned features.

# XOR: why hidden features help



A hidden layer can transform the space.

Then a simple output layer can solve the task.



**An MLP is a differentiable program**

$$h = g(W_1x + b_1)$$

$$\hat{y} = \text{softmax}(W_2h + b_2)$$


Every operation is differentiable, so L17's training loop applies unchanged.

# The code barely changes

```
model = nn.Sequential(  
    nn.Linear(d_in, 64),  
    nn.ReLU(),  
    nn.Linear(64, n_classes),  
)  
  
loss = loss_fn(model(x), y)  
loss.backward()  
optimizer.step()
```

The function is richer; the training pattern is the same.

# Backpropagation assigns credit




# What do hidden layers learn?



Deep models compose simple features into useful abstractions.

# CNNs learn visual features



**Local filters** detect patterns in patches.

**Weight sharing** detects the same pattern anywhere.

Later: VLMs use image encoders to turn images into vectors/tokens.

# Word IDs have no geometry

A raw ID is arbitrary:

cat = 17, dog = 932

An embedding gives each item a vector:

$$\text{Embed}[\text{cat}] \in \mathbb{R}^d$$

Now similarity can live in vector space.

# An embedding table is just parameters

| item | learned vector            |
|------|---------------------------|
| cat  | [0.22, -0.71, 0.10, ...]  |
| dog  | [0.19, -0.66, 0.12, ...]  |
| car  | [-0.80, 0.31, -0.40, ...] |

```
emb = nn.Embedding(  
    num_embeddings=vocab_size,  
    embedding_dim=128,  
)
```

Gradient descent updates these vectors during training.

# How does an embedding learn?

the cat sat on the \_\_\_\_

If the model should predict **mat**, gradients update the embeddings that helped make the prediction.

cat   dog      car   truck

geometry emerges from the training objective





# Embedding geometry is useful, not magic

## useful


nearest neighbors, retrieval,  
clustering

## dangerous

2D plots can overstate  
structure

Distances mean what the training objective made them mean.

# Inspect embeddings by projecting them




**PCA:** linear view of high-dimensional vectors.

**t-SNE/UMAP:** nonlinear visualization.

Use plots to ask questions, not as proof.

# Cluster embeddings with k-means



After embeddings exist, k-means becomes practical.


- inspect datasets
- find groups or duplicates
- build codebook intuition

# Autoencoders learn compression



Latent diffusion will generate in a compressed latent space.

# Contrastive learning aligns modalities



Matched image/text pairs are pulled together.

Mismatched pairs are pushed apart.

This is the bridge to VLMs and retrieval.

# Why this matters for LLMs

## **tokens**

become embeddings

## **attention**

moves information  
between  
embeddings

## **transformers**

scale this up

# Recap

- ✓ Neural nets learn **features**, not just predictions.
- ✓ Backprop assigns credit through the graph.
- ✓ **Embeddings** are learned vectors.
- ✓ PCA, t-SNE, and k-means help inspect embeddings.
- ✓ Autoencoders and CLIP prepare us for diffusion and VLMs.

## Next question

A sentence is a sequence of embeddings.

How should each word look at the other words?

L19: sequence models, attention, and transformers.