

# CS 486/686

# Decision Trees

Yuntian Deng

---

Lecture 18

RN 19.3 · PM 7.3.1

Search ›

Uncertainty ›

Decisions ›

Learning


## Learning goals

- Describe the **components** of a decision tree.
- **Construct** a decision tree given a feature-testing order.
- Compute a tree's **prediction accuracy** on a test set.
- Compute the **entropy** of a distribution.
- Compute the **expected information gain** of a feature.
- Trace + implement the **DT learning algorithm**.


# Decision trees in the wild

A decision tree is nothing more than a **nested if-then-else** — learned automatically from data.


## Should I use an emoji?



## What pet?



## Which language?



Simple to read, fast to evaluate, and surprisingly effective in practice.

# Running example: Jeeves the valet

Jeeves wants to predict whether Bertie will play tennis today. He has been recording the weather and Bertie's behavior every morning for 14 days.

## Features and target

- **Input features:** Outlook (Sunny/Overcast/Rain), Temp (Hot/Mild/Cool), Humidity (High/Normal), Wind (Weak/Strong).
- **Target:** Tennis? (Yes / No).


14-day training set + 14-day held-out test set. Goal: build a classifier from training, evaluate on test.

# The Jeeves training set

| Day | Outlook  | Temp | Humidity | Wind   | Tennis? |
|-----|----------|------|----------|--------|---------|
| 1   | Sunny    | Hot  | High     | Weak   | No      |
| 2   | Sunny    | Hot  | High     | Strong | No      |
| 3   | Overcast | Hot  | High     | Weak   | Yes     |
| 4   | Rain     | Mild | High     | Weak   | Yes     |
| 5   | Rain     | Cool | Normal   | Weak   | Yes     |
| 6   | Rain     | Cool | Normal   | Strong | No      |
| 7   | Overcast | Cool | Normal   | Strong | Yes     |
| 8   | Sunny    | Mild | High     | Weak   | No      |
| 9   | Sunny    | Cool | Normal   | Weak   | Yes     |
| 10  | Rain     | Mild | Normal   | Weak   | Yes     |
| 11  | Sunny    | Mild | Normal   | Strong | Yes     |
| 12  | Overcast | Mild | High     | Strong | Yes     |
| 13  | Overcast | Hot  | Normal   | Weak   | Yes     |

# Anatomy of a decision tree

- A supervised classifier with a **single discrete target**.
- Each **internal node** tests an input feature.
- **Edges** are labeled with feature values.
- Each **leaf** predicts a target class.
- **Classify**: follow edges from the root down to a leaf.



Sketch only — we'll build the real tree shortly.

# Classify an example: walk down the tree

Given a new example, follow the tests from the root to a leaf:

Example: **Outlook = Overcast, Temp = Hot, Humidity = High, Wind = Weak.**

Test Outlook → value is Overcast → leaf **Yes**. Done.

**Observation:** the path itself is a nested if-then-else program. Converting a tree to code is straightforward.

# Build a tree given an order

Pick a feature-testing order; recursively split the training examples by that feature's value.


## An order for the Jeeves data

- First, test **Outlook**.
- For Outlook = Sunny, test **Temp**.
- For Outlook = Rain, test **Wind**.
- (For other branches, test Humidity before Wind.)

Each example flows down the tree according to its feature values; leaves end up with subsets of the training set.



# The Jeeves decision tree



Built from the 14 training rows by following the feature order above.

# When do we stop? Three cases

## All same class

Every example at this node has the same label.

Return that label.

## No features left

We've tested every feature, but examples still disagree (noisy data).

Return the **majority class** at this node.

## No examples left

A feature value never appeared in training data — no examples reach this leaf.

Return the **majority class at the parent node**.

"No features left" handles noisy labels; "no examples left" handles unseen feature combinations.

# The decision-tree learner

## DT-Learner(*examples*, *features*)


1. If all examples are in the same class, return that class.
2. Else if no features left, return the majority class of *examples*.
3. Else if no examples left, return the majority class at the *parent*.
4. Else: choose a feature  $f$ ; for each value  $v$  of  $f$ ,
  - build an edge labeled  $v$ ;
  - recurse on examples with  $f = v$  and features  $\setminus \{f\}$ .

The only open question: **which feature do we pick to test?**

# Which feature should we test?

Pick the feature whose split makes the children most **class-pure**.

## Split on Outlook



Overcast cleanly predicts Yes — one

## Split on Temp



Every Temp value still mixes Yes/No — messier.


# Entropy: measuring uncertainty

## Entropy

$$I(P(c_1), \dots, P(c_k)) = - \sum_{i=1}^k P(c_i) \log_2 P(c_i)$$

Bits of uncertainty. Bigger = less certain.

- $I(0.5, 0.5) = 1$ : maximally uncertain.
- $I(0.01, 0.99) \approx 0.08$ : nearly certain the second outcome.
- $I(1, 0) = I(0, 1) = 0$ : no uncertainty at all.



# Expected information gain

If we split on a feature with  $k$  values, weighted-average the children's entropies:

## The three formulas

$$H_{\text{before}} = I\left(\frac{p}{p+n}, \frac{n}{p+n}\right)$$
$$H_{\text{after}} = \sum_{i=1}^k \frac{p_i + n_i}{p + n} I\left(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i}\right)$$
$$IG = H_{\text{before}} - H_{\text{after}}$$

Pick the feature with the **largest** information gain at every node.

# Jeeves: gain at the root (Outlook)

Training set has 9 Yes, 5 No ( $p = 9$ ,  $n = 5$ ).

**Q3.**  $H_{\text{before}} = I(9/14, 5/14) = ?$   
 $\Rightarrow$  **0.940 bits.**

**Q4.** Split on **Outlook**: Sunny (2+/3-), Overcast (4+/0-), Rain (3+/2-).

$$H_{\text{after}} = \frac{5}{14}(0.971) + \frac{4}{14}(0) + \frac{5}{14}(0.971) = 0.694.$$
$$\Rightarrow IG(\text{Outlook}) = 0.940 - 0.694 = \mathbf{0.247}.$$

# Jeeves: gain at the root (Humidity) and winner

Same  $H_{\text{before}} = 0.940$ . Now try splitting on Humidity instead.

**Q5.** Split on **Humidity**: High (3+/4-), Normal (6+/1-).

$$H_{\text{after}} = \frac{7}{14}(0.985) + \frac{7}{14}(0.591) = 0.788.$$

$$\Rightarrow IG(\text{Humidity}) = 0.940 - 0.788 = \mathbf{0.151}.$$

**Q6.** Which feature do we pick as the root?

**Outlook** – higher IG:  $0.247 > 0.151$ .



# ID3: the canonical algorithm

## **ID3(examples, features)**

1. Apply the three base cases from DT-Learner (same-class / no-features / no-examples).
2. For each feature  $f$  in features, compute  $IG(f)$  on the current examples.
3. Choose  $f^* = \arg \max_f IG(f)$ ; make an internal node testing  $f^*$ .
4. For each value  $v$  of  $f^*$ : recurse on examples with  $f^* = v$  and features  $\setminus \{f^*\}$ .

Greedy: pick the locally best feature at each step. Not always globally optimal — but fast and effective.

# Learning goals (recap) — Next: neural networks

- ✓ Describe the **components** of a decision tree.
- ✓ **Construct** a tree given a feature-testing order.
- ✓ Compute **test-set accuracy**.
- ✓ Compute the **entropy** of a distribution.
- ✓ Compute the **expected information gain**.
- ✓ Trace + implement the **DT learning algorithm**.

L19: from hand-engineered trees to **neural networks** — the building block of deep learning.