

# CS 486/686

## Neural Networks III

Yuntian Deng

---

Lecture 21

Optimization for deep learning · slides adapted from CMSC 35246 (Shubendu & Risi, U. Chicago)

Search ›

Uncertainty ›

Decisions ›


Learning

## Learning goals

- **Stochastic gradient descent:** noisy steps, learning-rate schedules, convergence.
- **Momentum** and its **Nesterov** look-ahead variant.
- **Adaptive learning rates:** AdaGrad and RMSProp.
- **Adam:** adaptive moments combining the two ideas.

# Roadmap: first-order optimizers

Backprop gives us  $\nabla_{\theta} L$ . The optimizer decides what to do with it.



- **SGD:** trade accuracy of gradient for speed of updates.
- **Momentum:** add inertia — escape ravines, damp oscillations.
- **Adaptive:** give each parameter its own learning rate.
- **Adam:** combine momentum + adaptive in one optimizer.

# SGD recap

**Batch GD:**

$$\hat{g} \leftarrow \frac{1}{N} \nabla_{\theta} \sum_i L(f(x^{(i)}; \theta), y^{(i)})$$

**SGD** (one example per step):

$$\hat{g} \leftarrow \nabla_{\theta} L(f(x^{(i)}; \theta), y^{(i)})$$

In both:  $\theta \leftarrow \theta - \varepsilon \hat{g}$ .

**Decay the learning rate linearly until iteration  $\tau$**   
:


$$\varepsilon_k = (1 - \alpha) \varepsilon_0 + \alpha \varepsilon_{\tau}, \quad \alpha = k/\tau$$

**Sufficient for convergence:**

$$\sum_k \varepsilon_k = \infty, \quad \sum_k \varepsilon_k^2 < \infty$$

**Mini-batch SGD** (e.g. 32, 64): cheaper than full-batch, less noisy than per-example, scales to huge datasets.

# What goes wrong with plain SGD



**Local minima.** Non-convex losses have many. SGD shrinks gradients to zero there and stalls.

**Noisy oscillations.** Each step uses only one example — the descent path zig-zags, slowing convergence.

Both issues motivate the next big idea:  
**momentum.**

# Momentum: add inertia


Introduce a **velocity** vector  $\mathbf{v}$  — an exponentially weighted average of past gradients:

$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \varepsilon \nabla_{\theta} L(\theta), \quad \theta \leftarrow \theta + \mathbf{v}$$

- $\alpha \in [0, 1)$  controls how much past gradients persist. Common:  $\alpha = 0.9$ .
- **Recent gradients dominate** — older ones decay geometrically.
- **Intuition:** a ball rolling down a hill keeps moving in the average direction, smoothing oscillations and powering through small bumps.

# Worked example

Minimize  $y = 0.3x^4 - 0.1x^3 - 2x^2 - 0.8x$ , with gradient  $g = 1.2x^3 - 0.3x^2 - 4x - 0.8$ . Start at  $x_0 = -2$ .



## SGD (no momentum)

| iter | $g_i$  | $\theta_i$ |
|------|--------|------------|
| 1    | -18.2  | -1.885     |
| 2    | -2.36  | -1.76      |
| 3    | -1.2   | -1.70      |
| 4    | -0.78  | -1.66      |
| ...  | ...    | ...        |
| 99   | 0.0002 | -1.586     |

stuck in local min

## With momentum ( $\alpha = 0.9$ )

| iter | $g_i$  | $v_i$  | $\theta_i$ |
|------|--------|--------|------------|
| 1    | -18.2  | 0      | -1.885     |
| 2    | -2.36  | -15.1  | -1.12      |
| 3    | 1.61   | -9.0   | -0.67      |
| 4    | 1.39   | -9.2   | -0.21      |
| ...  | ...    | ...    | ...        |
| 99   | 0.0002 | 0.0003 | 2.042      |

reaches global min

Velocity carries the iterate over the local-min barrier near  $x \approx -1.6$ .

# How big is one momentum step?

Plain SGD takes a step of size  $\varepsilon \|g\|$ . With momentum, recent gradients compound:

$$\text{step} \approx \varepsilon \|g_1\| + \alpha \varepsilon \|g_2\| + \alpha^2 \varepsilon \|g_3\| + \dots \approx \frac{\varepsilon \|\hat{g}\|}{1 - \alpha}$$

**Concretely:**  $\alpha = 0.9 \Rightarrow$  effective step size is up to **10 $\times$**  the per-step SGD step in a consistent direction.




In oscillating directions, opposite gradients cancel in  $\mathbf{v}$  — momentum naturally damps noise.



# Nesterov momentum: look ahead first

Idea: step in the direction of accumulated momentum, then evaluate the gradient there as a *correction*.



**Standard momentum:**

$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \varepsilon \nabla_{\theta} L(\theta)$$

**Nesterov momentum:**


$$\mathbf{v} \leftarrow \alpha \mathbf{v} - \varepsilon \nabla_{\theta} L(\theta + \alpha \mathbf{v})$$

$$\theta \leftarrow \theta + \mathbf{v}$$


Gradient evaluated at the look-ahead point gives a better local correction.

# Why one learning rate isn't enough

Features differ in scale and frequency. A single  $\varepsilon$  is too small for flat directions, too large for steep ones.



isotropic (easy)



anisotropic (zig-zag)

Give each parameter its **own** learning rate, scaled by its gradient history.

# AdaGrad: accumulate squared gradients

Per-parameter learning rate, scaled by the running sum of squared gradients.

## AdaGrad

*Require: global LR  $\varepsilon$ , small  $\delta$ ; init  $\mathbf{r} = \mathbf{0}$ .*

1. Sample example  $(x^{(i)}, y^{(i)})$ ; compute  $\hat{\mathbf{g}} \leftarrow \nabla_{\theta} L(f(x^{(i)}; \theta), y^{(i)})$ .
2. Accumulate:  $\mathbf{r} \leftarrow \mathbf{r} + \hat{\mathbf{g}} \odot \hat{\mathbf{g}}$
3. Update:  $\Delta\theta \leftarrow -\frac{\varepsilon}{\delta + \sqrt{\mathbf{r}}} \odot \hat{\mathbf{g}}$
4. Apply:  $\theta \leftarrow \theta + \Delta\theta$

**Catch:**  $\mathbf{r}$  only grows — the effective LR shrinks to zero, training stalls in non-convex settings.

# RMSProp: forget old gradients

Fix AdaGrad's shrink by replacing the running sum with an **exponentially decaying** moving average.

## RMSProp

*Require: LR  $\varepsilon$ , decay  $\rho$  (typically 0.9), small  $\delta$ ; init  $\mathbf{r} = 0$ .*

1. Compute gradient:  $\hat{g} \leftarrow \nabla_{\theta} L(f(x^{(i)}; \theta), y^{(i)})$ .
2. Decay+accumulate:  $\mathbf{r} \leftarrow \rho \mathbf{r} + (1 - \rho) \hat{g} \odot \hat{g}$
3. Update:  $\Delta\theta \leftarrow -\frac{\varepsilon}{\delta + \sqrt{\mathbf{r}}} \odot \hat{g}$
4. Apply:  $\theta \leftarrow \theta + \Delta\theta$

Related: **AdaDelta** goes one step further and removes the global LR  $\varepsilon$  by also tracking  $\mathbb{E}[\Delta\theta^2]$ .

# Adam: ADAdaptive Moments

Track **first moment** (mean of  $\hat{g}$ ) like momentum, and **second moment** (mean of  $\hat{g}^2$ ) like RMSProp.

## Adam

Require: LR  $\varepsilon$ , decay rates  $\rho_1, \rho_2$  (typically 0.9, 0.999), small  $\delta$ ; init  $\mathbf{s} = \mathbf{0}, \mathbf{r} = \mathbf{0}, t = 0$ .

1. Compute gradient  $\hat{g}$ ; increment  $t$ .
2. **1st moment**:  $\mathbf{s} \leftarrow \rho_1 \mathbf{s} + (1 - \rho_1) \hat{g}$
3. **2nd moment**:  $\mathbf{r} \leftarrow \rho_2 \mathbf{r} + (1 - \rho_2) \hat{g} \odot \hat{g}$
4. **Bias correction**:  $\hat{\mathbf{s}} \leftarrow \frac{\mathbf{s}}{1 - \rho_1^t}, \quad \hat{\mathbf{r}} \leftarrow \frac{\mathbf{r}}{1 - \rho_2^t}$
5. Update:  $\Delta\theta = -\varepsilon \frac{\hat{\mathbf{s}}}{\sqrt{\hat{\mathbf{r}} + \delta}}$
6. Apply:  $\theta \leftarrow \theta + \Delta\theta$

Bias correction fixes the cold-start: early  $\mathbf{s}, \mathbf{r}$  are biased toward 0 because they start at 0.

# Optimizer cheat sheet

| Method     | Update direction                                                         | Per-param LR? | When to use                               |
|------------|--------------------------------------------------------------------------|---------------|-------------------------------------------|
| SGD        | $-\varepsilon \hat{g}$                                                   | No            | Simple losses; baseline.                  |
| + Momentum | $\mathbf{v} = \alpha \mathbf{v} - \varepsilon \hat{g}$                   | No            | Noisy gradients, ravines, local minima.   |
| RMSProp    | $-\frac{\varepsilon}{\sqrt{\mathbf{r}} + \delta} \hat{g}$                | Yes           | RNNs, non-stationary or sparse gradients. |
| Adam       | $-\varepsilon \frac{\hat{\mathbf{s}}}{\sqrt{\hat{\mathbf{r}}} + \delta}$ | Yes           | Default choice in deep learning today.    |

Tune the learning rate first; pick the optimizer second.

# Revisiting learning goals

- **Stochastic gradient descent:** noisy steps, learning-rate schedules, convergence.
- **Momentum** and its **Nesterov** look-ahead variant.
- **Adaptive learning rates:** AdaGrad and RMSProp.
- **Adam:** adaptive moments combining the two ideas.

Search ›

Uncertainty ›

Decisions ›

Learning ›

Recap

## Next: L22 — course recap & exam prep

- A tour of everything we learned: search, CSPs, probabilistic reasoning, decision theory, MDPs, RL, ML.
- How the pieces fit into **modern AI systems**.
- Pointers to follow-up courses (CS480/680, CS485) and open research.
- **Final exam logistics:** what to bring, format, and a notation cheat sheet for your study card.

Thanks for a great term — see you on Tuesday!