

# CS 486/686

## Course Recap

Yuntian Deng

---

Lecture 22

Putting the term together: search, uncertainty, decisions, learning

# The course in one picture



# Search I – uninformed

## Problem formulation

- States, actions, transition model, initial & goal states, path cost.
- Search tree explores states by expanding the **frontier**.

*Pick the data structure for the frontier → pick the algorithm.*

## DFS vs BFS vs IDS

- **DFS**: stack →  $O(bm)$  space, not complete.
- **BFS**: queue → complete,  $O(b^d)$  time and space.
- **IDS**: DFS in increasing depth → BFS-completeness, DFS-space.

*IDS combines the best of both.*

# Search II — heuristic + A\*

## The three frontier rules

- **LCFS:**  $\min \text{cost}(n)$ .
- **GBFS:**  $\min h(n)$ .
- **A\*:**  $\min f(n) = \text{cost}(n) + h(n)$ .

## Heuristic properties

- **Admissible** ( $h \leq h^*$ )  $\Rightarrow$  A\* is optimal with multi-path pruning off.
- **Consistent** ( $h(n) \leq c(n, n') + h(n')$ )  $\Rightarrow$  safe to use multi-path pruning.
- Consistent  $\Rightarrow$  admissible (not the other way around).

# Search strategy summary

| Strategy          | Frontier choice              | Complete? | Space  | Time |
|-------------------|------------------------------|-----------|--------|------|
| Depth-first       | Last added                   | No        | Linear | Exp  |
| Breadth-first     | First added                  | Yes       | Exp    | Exp  |
| Lowest-cost first | $\min \text{cost}(n)$        | Yes       | Exp    | Exp  |
| Greedy best-first | $\min h(n)$                  | No        | Exp    | Exp  |
| A*                | $\min \text{cost}(n) + h(n)$ | Yes       | Exp    | Exp  |

A\* dominates when the heuristic is informative enough to prune most of the tree.

# CSPs + local search

## Constraint satisfaction

- Generate-and-test is exponential — exploit problem structure.
- **Backtracking**: incremental, prune as soon as a constraint fails.
- **Arc consistency / AC-3**: shrink domains before/during search; complexity  $O(cd^3)$ .

## Local search

- Trade completeness for very large state spaces.
- **Greedy descent / hill-climbing**: always move to best neighbour.
- Beats local minima with **random restarts** + **simulated annealing** (accept worse moves with prob  $e^{-\Delta/T}$ ).

# Probability + independence

## The four rules

- **Product:**  $P(A, B) = P(A \mid B) P(B)$
- **Sum:**  $P(A) = \sum_b P(A, B=b)$
- **Chain:**  $P(X_1, \dots, X_n) = \prod_i P(X_i \mid X_{<i})$
- **Bayes:**  $P(H \mid e) \propto P(e \mid H) P(H)$

## Independence

- **Unconditional:**  $P(A, B) = P(A) P(B)$ .
- **Conditional:**  $P(A, B \mid C) = P(A \mid C) P(B \mid C)$ .
- Tells us which terms drop out of the chain rule — the key to scaling up inference.

# Bayesian networks + d-separation

## Construction recipe

- Pick a variable order (causal often works best).
- Add each node; choose the smallest set of parents that makes it conditionally independent of the rest.
- Joint factorizes as  $P(X_1, \dots, X_n) = \prod_i P(X_i \mid \text{pa}(X_i))$ .

## 3 d-separation patterns

 **Chain**  $A \rightarrow B \rightarrow C$ : blocked by observing  $B$ .

 **Common cause**  $A \leftarrow B \rightarrow C$ : blocked by observing  $B$ .

 **Common effect**  $A \rightarrow B \leftarrow C$ : blocked by *not* observing  $B$  or descendants.

*A path is blocked  $\Rightarrow$  the endpoints are conditionally independent.*



# Inference — VE & HMM filtering

## Variable elimination

- **Define** a factor for each CPT.
- **Restrict** factors to evidence.
- **Multiply** factors sharing a variable; **sum out** the hidden one.
- **Normalize** to get a posterior.


## HMM filtering

- Recursion:  $P(S_t \mid o_{0:t}) \propto P(o_t \mid S_t) \sum_{s_{t-1}} P(S_t \mid s_{t-1}) P(s_{t-1} \mid o_{0:t-1})$
- One forward pass:  $O(t \cdot |S|^2)$  time.

# HMM decoding – smoothing & Viterbi

## Forward×backward smoothing

- Forward:  $\alpha_k = P(S_k \mid o_{0:k})$ .
- Backward:  $\beta_k = P(o_{k+1:t-1} \mid S_k)$ .
- Posterior:  $P(S_k \mid o_{0:t-1}) \propto \alpha_k \beta_k$ .



# Viterbi

- DP over the trellis, max instead of sum.
- Backtrack the argmax pointers to recover the MAP sequence.

# Decision networks

## Three node types

- **Chance** nodes — like a Bayes net (CPTs).
- **Decision** nodes — rectangles, no CPT, chosen by us.
- **Utility** node — diamond, deterministic function of its parents.

## Maximum expected utility

- Policy  $\pi(d \mid \text{parents})$ ; MEU picks  $\pi^*$  maximizing  $\mathbb{E}[U]$ .
- Evaluate by enumeration or by running VE on the DN.
- **Value of information:** EU gain from observing a variable before deciding.

# Markov decision processes

## Setup

- States, actions, transition  $P(s' \mid s, a)$ , reward  $R(s)$ , discount  $\gamma$ .
- Policy  $\pi : S \rightarrow A$ ; value  $V^\pi(s) = \mathbb{E}[\sum_t \gamma^t R(s_t) \mid \pi]$ .

## Bellman + the two algorithms

- **Bellman optimality:**  $V^*(s) = R(s) + \gamma \max_a \sum_{s'} P(s' \mid s, a) V^*(s')$ .
- **Value iteration:** apply the Bellman update until convergence.
- **Policy iteration:** evaluate (linear system)  
↔ improve (greedy) — converges fast.

# Reinforcement learning

## Two learning regimes

- **Passive ADP:** follow a fixed policy, estimate  $\hat{P}, \hat{R}$ , solve via PE.
- **Active ADP:** also pick actions — balance exploration vs exploitation ( $\epsilon$ -greedy, softmax, optimistic init).

## Model-free updates

- **Q-learning** (off-policy):  $Q(s, a) \leftarrow Q(s, a) + \alpha [R + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ .
- **SARSA** (on-policy): use the action actually taken instead of  $\max_{a'}$ .
- ADP converges fast but needs a model; Q-learning is slower but model-free.

# Supervised + unsupervised learning

## Supervised

- **Classification** (discrete  $y$ ) vs **regression** (continuous  $y$ ).
- **Bias-variance**: underfit (high bias)  $\leftrightarrow$  overfit (high variance).
- **Cross-validation**:  $K$ -fold to pick the model complexity that minimizes held-out error.
- **ERM**: minimize empirical loss  $\frac{1}{N} \sum_i \ell(\hat{y}^{(i)}, y^{(i)})$ .

## Unsupervised

- **k-means**: alternate assign  $\leftrightarrow$  recompute centroids; pick  $k$  via elbow / silhouette.
- **PCA**: project onto top eigenvectors of the covariance matrix; maximize variance.
- **Autoencoders & GANs**: learn representations / generators with neural nets.

# Decision trees + neural networks

## Decision trees

- Greedy top-down split with **ID3**.
- Choose split that maximizes **information gain**  $I = H(Y) - H(Y | X)$ .
- Stop when pure / no features / no examples.
- Interpretable, great for small tabular data.

## Neural networks

- Weighted sums + non-linear activations (sigmoid, ReLU).
- **Forward pass** computes loss; **backprop** computes gradients via chain rule ( $\delta$  recursion).
- Train with **SGD** + momentum / Adam.
- Black box but unmatched for images, audio, text.

# Notation cheat sheet

Symbols that recur across the four modules. Watch for the two collisions on the right.

| Symbol          | Meaning                                                                                                     | First seen              |
|-----------------|-------------------------------------------------------------------------------------------------------------|-------------------------|
| $\gamma$        | Discount factor for future rewards                                                                          | L13 (MDPs)              |
| $\pi, \pi^*$    | Policy / optimal policy                                                                                     | L12–L15                 |
| $V(s), Q(s, a)$ | State / state-action value functions                                                                        | L13–L15                 |
| $H(p), IG$      | Entropy / information gain ( $IG = H_{\text{before}} - H_{\text{after}}$ )                                  | L6 (entropy) · L18 (IG) |
| $\mathcal{H}$   | Hypothesis class (ERM)                                                                                      | L16                     |
| $\Sigma$        | Covariance matrix (PCA)                                                                                     | L17                     |
| $\alpha$        | RL TD step size (L15); momentum coefficient (L21) — <i>different meanings</i>                               | L15 / L21               |
| $\epsilon$      | Exploration probability in $\epsilon$ -greedy (L15); learning rate in SGD (L21) — <i>different meanings</i> | L15 / L21               |

Worth a row on your two-sided study sheet.



# Final exam — logistics

## What to bring

- Non-programmable calculator.
- One two-sided A4 study sheet (handwritten or typed).

## Format (tentative)

- Around 8 problems spanning all four modules.
- ~84 marks total; 76+ marks count as 100%.
- Pass threshold: 42 marks.

*Exact details will be confirmed on the course page before exam day.*

Thanks for a great term — good luck on the final!