

CS 486/686

Dissecting Qwen3.5-0.8B

Yuntian Deng

Lecture 23

How a language model sees

Search ›

Uncertainty ›


Decisions ›

Learning

By the end, you can trace how Qwen sees

1. **Run** visual Q&A and OCR in your browser.
2. **Trace** pixels → patches → image embeddings.
3. **Calculate** how resolution changes the token count.
4. **Explain** how a familiar embedding interface lets us reuse an LLM.

Without an image, Qwen is the LM you already know



L21 already explained this path. Today we change the input, not the output loop.

The exact model we will dissect

Qwen3.5-0.8B

post-trained checkpoint

Qwen/Qwen3.5-0.8B

1024

LM hidden width

24

language blocks

16×16

image patches

text

output modality

Qwen Team, “Qwen3.5: Towards Native Multimodal Agents,” 2026; official model configuration.

One twist: most blocks carry a recurrent state




Long image contexts do not pay full quadratic attention in every block.

Ask Qwen about an image

LIVE IN BROWSER

Choose the photo—or upload your own—then ask a visual question.



RECORDED QWEN3.5-0.8B OUTPUT

QUESTION

Is the person wearing a suit? What are they wearing?

ANSWER

Yes, the person is wearing a suit. They are dressed in a dark suit jacket, a white collared shirt, and a dark tie.


Live path: Transformers.js v4 + WebGPU + onnx-community/Qwen3.5-0.8B-ONNX. No API key; inference stays in the browser.

THE CENTRAL QUESTION

The language model did not receive pixels.

What vectors did it receive?

Make the image look like language to the LLM



Pixels and token IDs begin in different worlds

image

[384, 640, 3]

737,280 RGB numbers


text

[2488, 374, 2201, ...]

a short sequence of IDs

First compress local pixels into a sequence of patch vectors.

Qwen starts with 16×16 image patches



one patch

$16 \times 16 \times 3 = 768$ pixel values

one learned vector

the patch embedding

all patches

a spatial sequence

Patchify while preserving the image shape LIVE


Change the retained resolution; count raw patches and merged image embeddings.

640×384 image → 40×24 raw patches → 20×12 = **240 image embeddings**

A fixed square resize can erase small details



original CLIP-style input
224×224 square



variable-resolution input
640×384, aspect preserved

Documents need enough retained patches for their text to survive.

“Native resolution” means a variable patch grid

preserve

aspect ratio

round

dimensions to patch-
compatible sizes

cap

the total visual-token budget

The processor does not promise to retain every original pixel;
it chooses the largest useful grid within its configured budget.

SigLIP 2 NaFlex and Qwen multimodal image preprocessing use variable grids to reduce aspect-ratio distortion.

More retained pixels create a longer visual sequence

256×256

16×16 = 256 raw patches

64 after 2×2 merge

512×512

32×32 = 1024 raw patches

256 after 2×2 merge


640×384

40×24 = 960 raw patches

240 after 2×2 merge

Resolution buys detail by spending context and compute.

Qwen merges each 2×2 neighborhood



$$N_{\text{LM image}} = \frac{H}{16} \frac{W}{16} \frac{1}{4} = \frac{H}{32} \frac{W}{32}$$

PATCHES ARE ONLY THE START

**How does a patch come to
mean “cup” or “deadline”?**

It must be trained to carry visual semantics.

Recall CLIP: put matching images and text nearby





image encoder

$$f_I(I)$$



*“a person wearing
a suit”*

text encoder

$$f_T(T)$$

Radford et al., “Learning Transferable Visual Models From Natural Language Supervision,” 2021. Callback to L18.

Each image finds its caption—and vice versa LIVE

The diagonal contains real pairs; other cells are in-batch mismatches.

photo ↔ “person”	photo ↔ “chart”
chart ↔ “person”	chart ↔ “chart”

Contrastive training turns appearance into meaning

before


colors, edges, textures

after matching text

objects, actions, concepts

Now the vision features are useful to a language model—not merely compact pixels.

Qwen's vision tower is derived from SigLIP2



We use the idea and move on; the SigLIP2 training refinements are optional reading.

Tschannen et al., “SigLIP 2,” 2025; Qwen3.5 vision configuration.

The vision tower contextualizes every patch

patch embeddings



12 vision blocks



$[N, 768]$ contextual features

HIDDEN WIDTH

768

ATTENTION HEADS


12

PATCH SIZE

16×16

The same tower also handles video with temporal patches; today we follow one still image.

The connector changes count and width



Qwen3.5 visual merger: LayerNorm, 2×2 spatial merge, two linear layers with GELU.

Same shape means the same LLM interface

image embeddings

[V, 1024]

from the connector

word embeddings


[T, 1024]

from the lookup table

```
llm(inputs_embeds = concat(image_embeds, text_embeds))
```

To the backbone, both are sequences of 1024-dimensional vectors.

Turn the unknown into the known




This pattern appears everywhere in computer science: adapt, reduce, reuse.

The image and question become one context




The image embeddings are context before the answer—just like earlier words.

Once embedded, generation is unchanged



End-to-end training teaches Qwen to use the adapter



One image, every boundary



The recorded trace is generated from the real processor;
the demo reports its actual grid for uploaded images.

The complete browser path is short

```
const processor = await AutoProcessor.from_pretrained(MODEL);
const model = await Qwen3_5ForConditionalGeneration
  .from_pretrained(MODEL, {
    device: "webgpu",
    dtype: {
      embed_tokens: "q4",
      vision_encoder: "fp16",
      decoder_model_merged: "q4",
    },
  });

const prompt = processor.apply_chat_template(messages, {
  add_generation_prompt: true,
  enable_thinking: false,
});
const inputs = await processor(prompt, [image]);
const output = await model.generate({ ...inputs, max_new_tokens: 80 });
```

Adapted from the official Transformers.js Qwen3.5 WebGPU example.

processor

patch + serialize

model

encode + adapt + reuse LM

generate

emit text tokens

Read a document with the same model LIVE IN BROWSER

The model loaded for VQA is reused—choose the announcement or upload a document.

CS 486/686
ADDITIONAL OFFICE HOURS

WED JUL 29	3-4 PM	DC 2633
THU JUL 30	4-5 PM	CLASSROOM

Bring questions from L19-L23.

RECORDED QWEN3.5-0.8B OUTPUT

QUESTION

Where and when are the Wednesday office hours?

ANSWER

Wednesday 3-4 PM at DC 2633

Too few patches can erase the evidence

CS 486/686 ADDITIONAL OFFICE HOURS		
WED JUL 29	3-4 PM	DC 2633
THU JUL 30	4-5 PM	CLASSROOM

Bring questions from L19-L23.

240 embeddings

Qwen: "Wednesday 3-4 PM at DC 2633"

CS 486/686 ADDITIONAL OFFICE HOURS		
WED JUL 29	3-4 PM	DC 2633
THU JUL 30	4-5 PM	CLASSROOM

Bring questions from L19-L23.

77 embeddings

Qwen: "Wednesday 3:00 PM" — room lost

The model cannot recover visual evidence that preprocessing removed.

Recap: adapt the image, then reuse the LLM

- 1 Patch** the image into 16×16 regions.

- 2 Encode** them as 768-wide semantic features.

- 3 Merge + project** 2×2 groups into 1024-wide image embeddings.

- 4 Concatenate** image and word embeddings through the same interface.

- 5 Reuse Qwen** to predict the text answer.

LOOK AT THE OUTPUT HEAD

Qwen can see images—but it can only emit text tokens.

